



Carnegie Mellon University
School of Computer Science

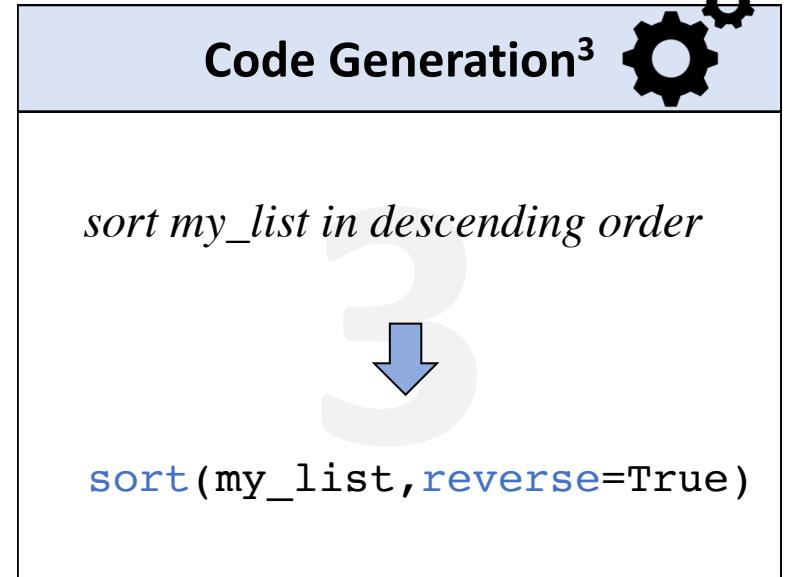
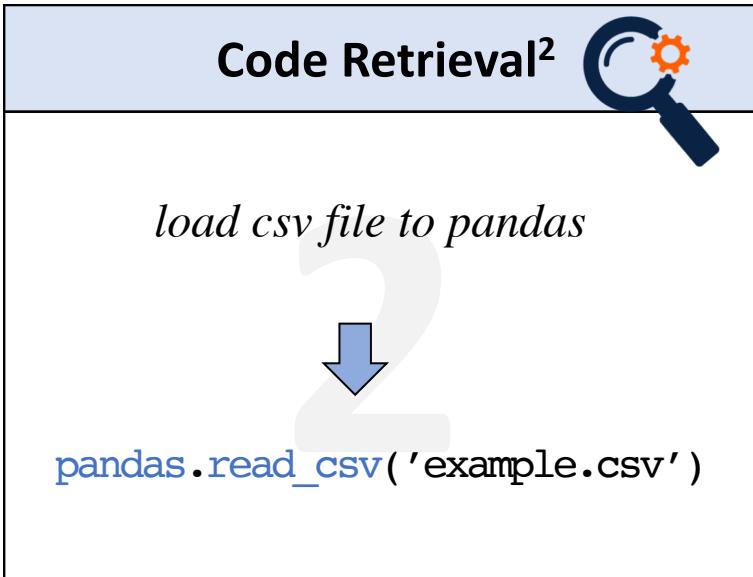
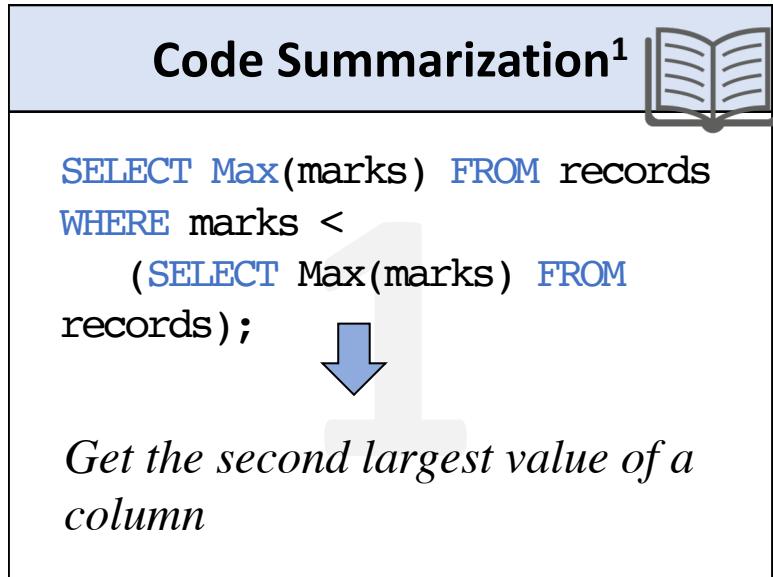
Learning to Mine Aligned Code and Natural Language Pairs from Stack Overflow

Pengcheng Yin* Bowen Deng* Edgar Chen Bogdan Vasilescu Graham Neubig

Carnegie Mellon University

Background

- When Natural Language Processing (NLP) meets Software Engineering...



- These tasks, mostly powered by **data-driven models**, heavily rely on parallel training (and evaluation) corpora of source code and natural language in **high quality** and **large amount**

- [1] Iyer et al., ACL 2016
- [2] Zhang et al., FSE 2016
- [3] Yin and Neubig., ACL 2017, ACL 2018



Collecting Intent/Snippet Pairs from SO

- Such data-driven models require parallel data of natural language **intents** and source code **snippets** and in high volume and high quality
 - Intent** natural language description of what a programmer would like to do
 - Snippet** a piece of source code that implements the intent

Intent get the maximum value of a column

Snippet `SELECT MAX(marks) from records`

Intent read a csv file into pandas

Snippet `pandas.read_csv('example.csv')`

Intent sort my_list in descending order

Snippet `sort(my_list, reverse=True)`

:

Can we collect such data from



Heuristic approaches [Wong et al., 2013; Iyer et al., 2016]?

- Select **all** code blocks
- Select **all** code blocks in **accepted answers**

Are these Heuristic Approaches Good Enough?

Intent

Removing duplicates in lists

▲
566 Pretty much I need to write a program to check if a list has any duplicates and if it does it removes them and returns a new list with the items that weren't duplicated/removed. This is what I have but to be honest I do not know what to do.

```
def remove_duplicates():
    t = ['a', 'b', 'c', 'd']
```

▲
1000 The common approach to get a unique collection of items is to use a `set`. Sets are *unordered* collections of *distinct* objects. To create a set from any iterable, you can simply pass it to the built-in `set()` function. If you later need a real list again, you can similarly pass the set to the `list()` function.

The following example should cover whatever you are trying to do:

```
>>> t = [1, 2, 3, 1, 2, 5, 6, 7, 8]
>>> t
[1, 2, 3, 1, 2, 5, 6, 7, 8]
>>> list(set(t))
[1, 2, 3, 5, 6, 7, 8]
>>> s = [1, 2, 3]
>>> list(set(t) - set(s))
[8, 5, 6, 7]
```

Contextual
Information
Actual Snippet
Auxiliary info.
Irrelevant Code

As you can see from the example result, the original order is not maintained. As mentioned above, sets themselves are unordered collections, so the order is lost. When converting a set back to a list, an arbitrary order is created.

However, acquiring such data from Stack Overflow posts may not be that straight-forward

1. **Contextual Information** package import statements, variable definition
2. **Auxiliary information** return values, example outputs
3. **Irrelevant Code**
4. Counter-examples

Heuristic approaches? 😞

- Select **all** code blocks
- Select **all** code blocks in **accepted answers**



Our Solution

- CONALA, a system to collect parallel data of source **code** snippet and **natural language** intents from Stack Overflow
 - *data-driven*: learn patterns of “good” and “bad” intent/snippet pairs from data (using neural networks)
 - *language-agnostic*: applicable to different programming languages (e.g., Python and Java)
 - *Scalable*: capable of applying to full-scale Stack Overflow data (collected ~600K intent/snippet pairs for Python)

Project Website: *conala-corpus.github.io*



System Architecture

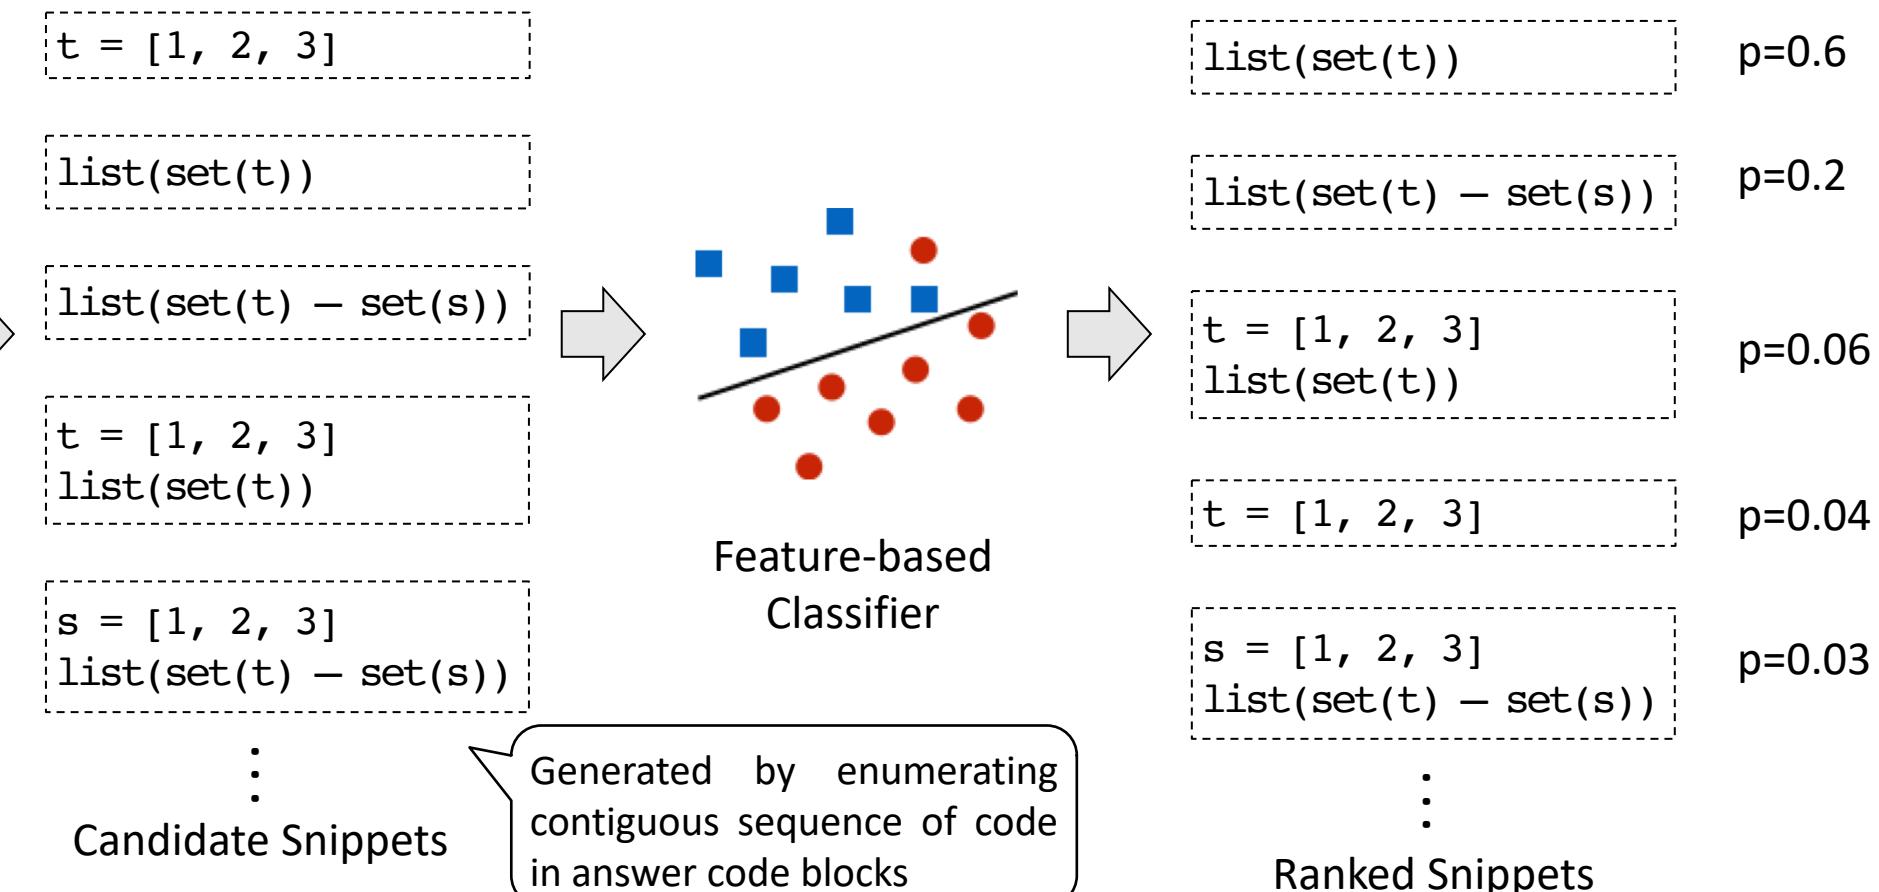
- Enumerate all possible candidate intent/snippet pairs in a Stack Overflow page
- Learn a classifier to rank each candidate intent/snippet pair

Question (i.e. intent)
removing duplicates in lists

An Answer Code Block

```
t = [1, 2, 3]
list(set(t))
s = [1, 2, 3]
list(set(t) - set(s))
```

Input SO question
("how-to" questions)



Features

- **Purpose** measure the probability (plausibility) of each intent/snippet candidate
- **Two types of features** language independent, highly-indicative

$$P\{correct \mid \begin{array}{l} \text{Intent} = \text{removing duplicates in lists} \\ \text{Snippet} = \text{list}(\text{set}(t)) \end{array}\}$$



Structural
Features



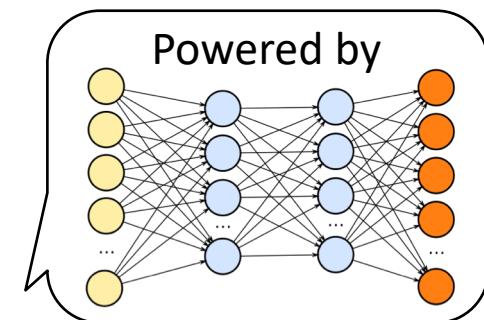
Correspondence
Features

Inspired by heuristic approaches, does **not** use intent information

Code Shape features: FULLBLOCK CONTAINSIMPORT STARTSWITHASSIGNMENT ...

Answer quality features: ISACCEPTEDANSWER POSTRANK ...

★ Use state-of-the-art **neural networks** to estimate the (semantic) correspondence between intent and snippet

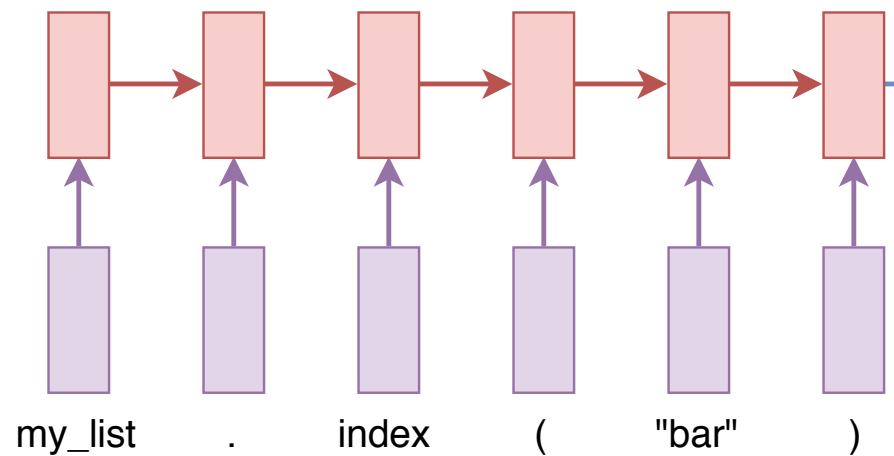


Score{ *removing duplicates in lists* \Leftrightarrow `list(set(t))` }

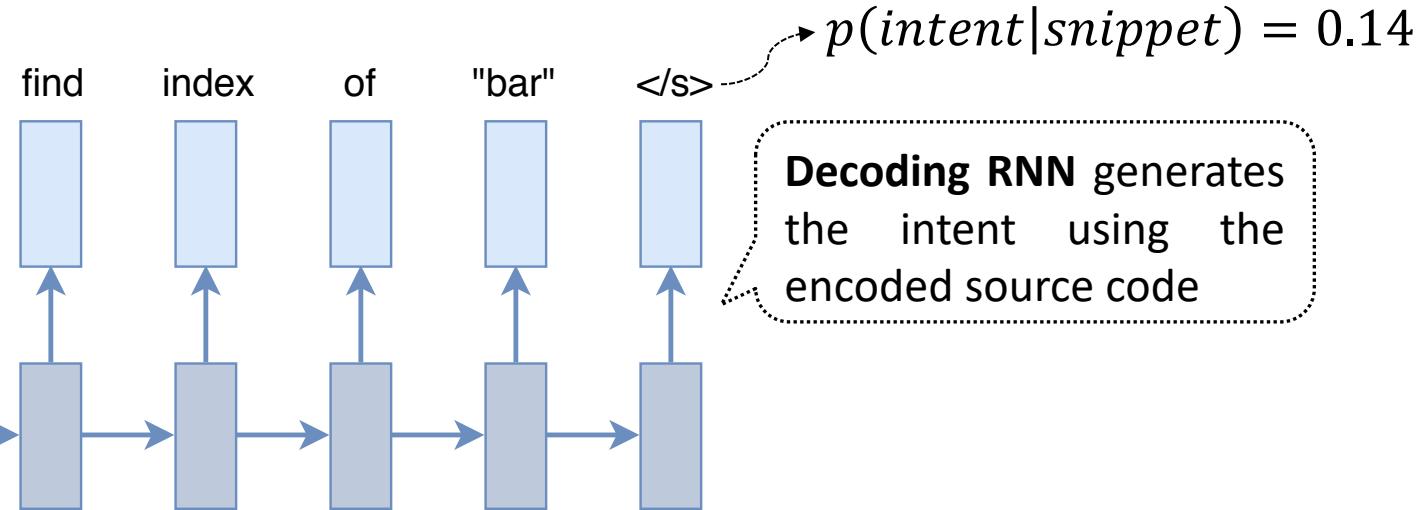
Neural Correspondence Model between Code and Intent

- **Neural sequence-to-sequence networks** for translation probability between intent and snippet

Encoding Recurrent Neural Network (RNN) encodes the information of the code



Neural Net for Snippet \mapsto Intent



Decoding RNN generates the intent using the encoded source code

No Need to Use Curated Training Data
using large amounts of question titles and
code blocks from accepted answers to train
the model 😊

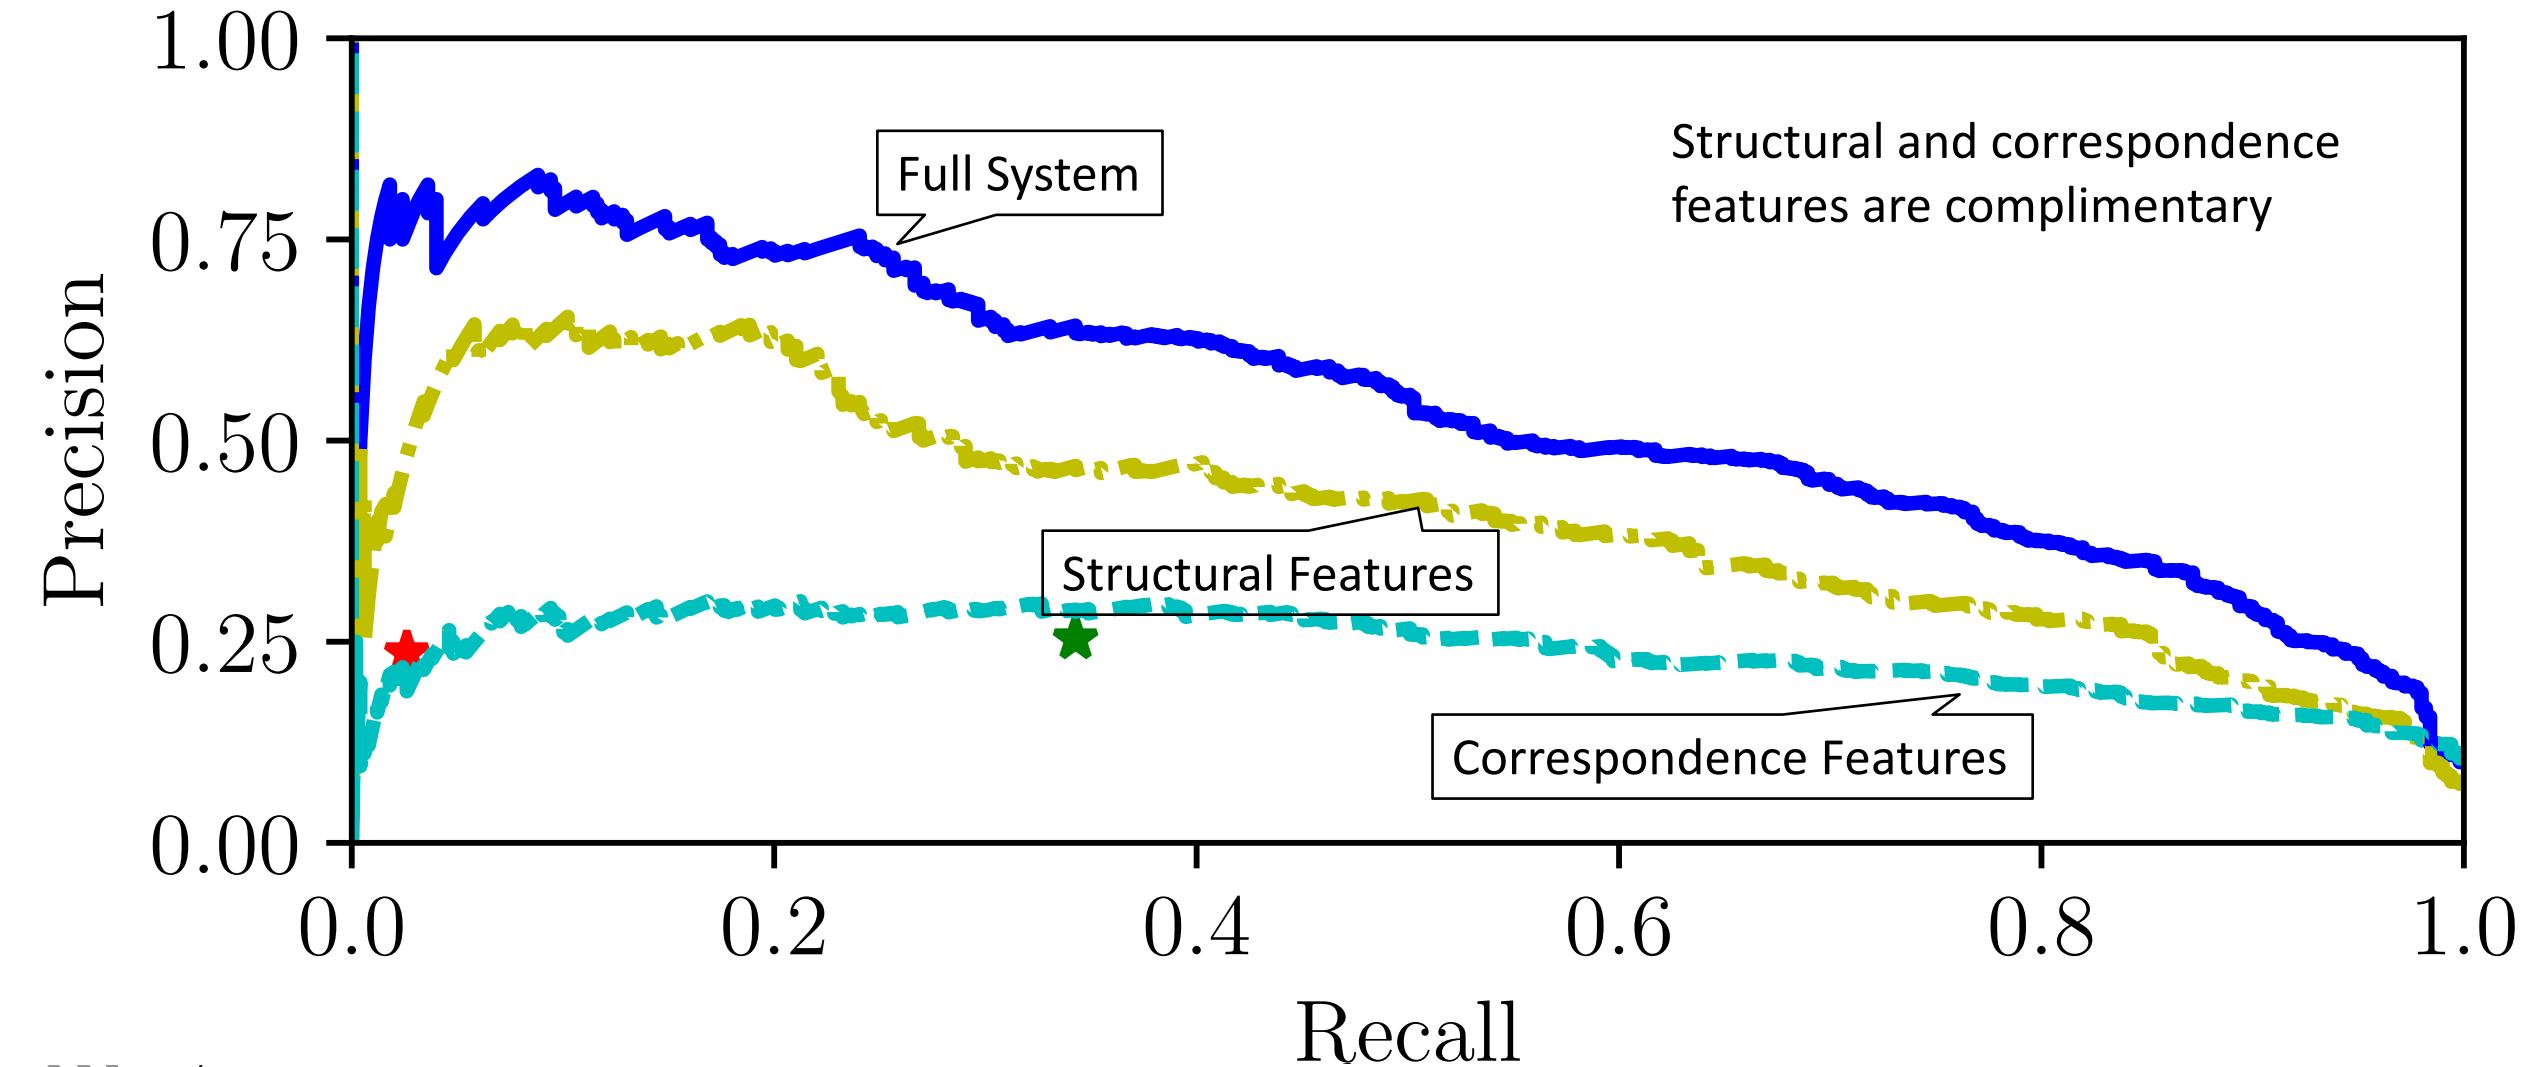
What Left to be Learned?

- **Structural Features** shown by existing works as indicative
- **Correspondence Features** already pre-trained on massive, readily available data on Stack Overflow (questions and code blocks)!
- We just need **small amount** of manually annotated intent/snippet pairs to tune the ~20 weights in the classifier 😊

	Gold Annotations
	527 examples 142 questions
	330 examples 100 questions

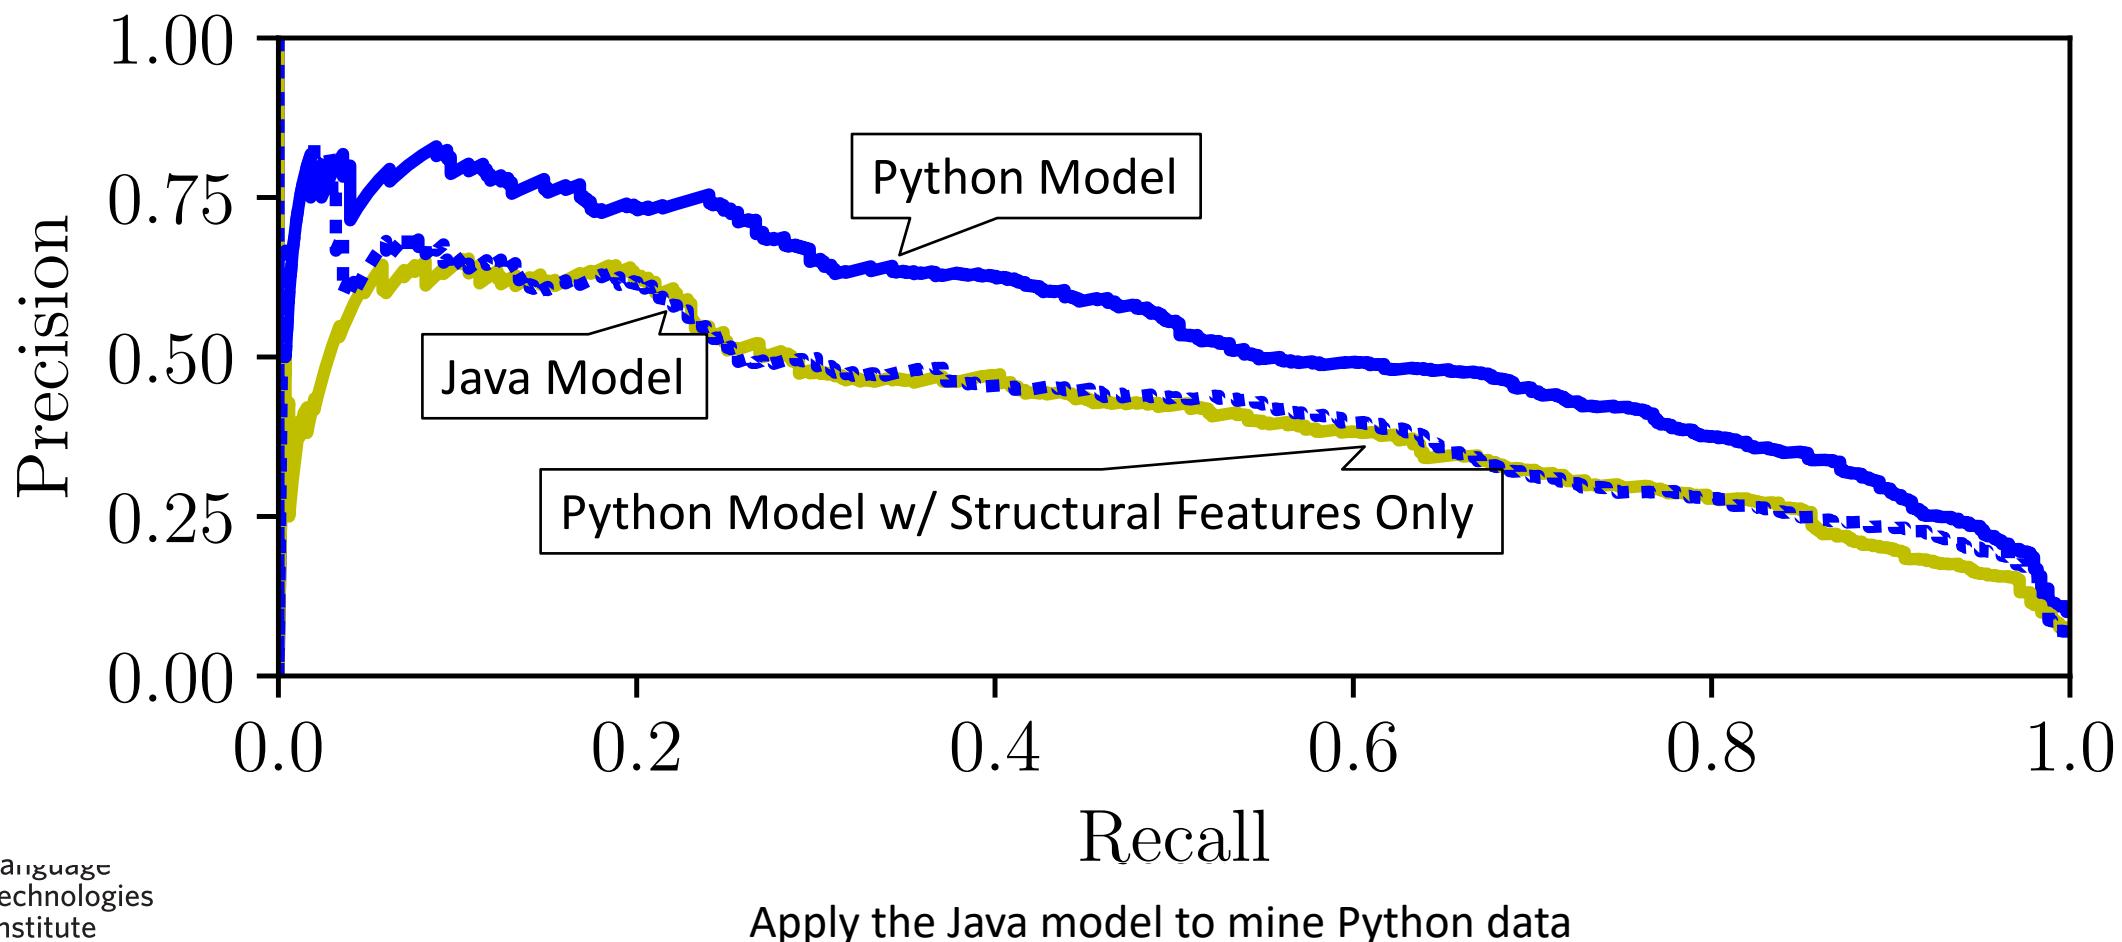


Cross-Validation Results



Must we Annotate each Language?

- The classifier requires small set of gold-standard annotations to train on
- When apply our system to a new language, can we use the existing model trained on another (old) language?



Dataset Collection

- Apply the system to Python questions on Stack Overflow, collecting ~600K pairs
- ~2300 (and counting) high-quality annotated intent/code snippet pairs
- **Rewritten Intents** manually annotated, revised intents to reflect the full meaning of the code
 - Add free variable names, arguments to the intent
 - Useful for fine-grained language to code tasks like code generation

Crowdsourced

Intent copying one file's contents to another

Rewritten Intent copy the content of file 'file.txt' to file 'file2.txt'

Code Snippet `shutil.copy('file.txt','file2.txt')`

An example from the annotation dataset



Examples

- Examples covers a wide variety of use cases
 - Built-in data type operation
 - I/O operation
 - Third-party library usage
- Examples are highly expressive and compositional!
 - Pose challenges to existing code/NL models

Intent dict how to create key or append an element to key

Rewritten Intent Create a key `key` if it does not exist in dict `dic` and append element `value` to value

Code Snippet `dic.setdefault(key, []).append(value)`

Intent How do I check if all elements in a list are the same

Rewritten Intent check if all elements in list `mylist` are the same

Code Snippet `len(set(mylist)) == 1`

Intent Iterate through words of a file in Python

Rewritten Intent check if all elements in list `mylist` are the same

Code Snippet `words = open('myfile').read().split()`

Intent Delete Column in Pandas based on Condition

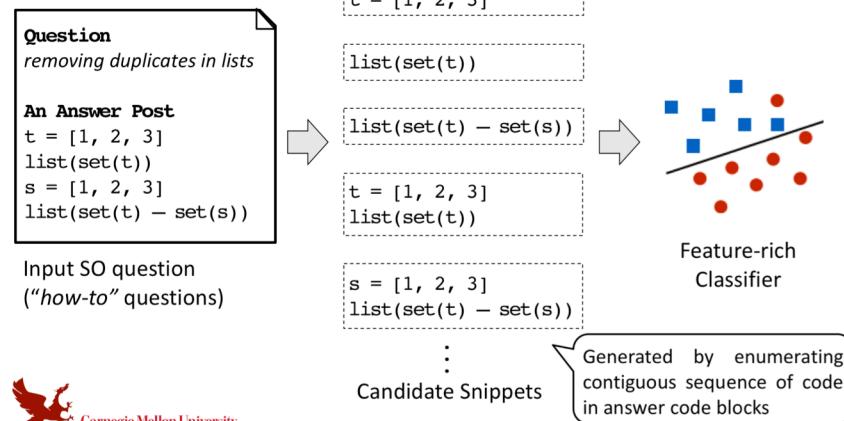
Rewritten Intent delete all columns in DataFrame `df` that do not hold a non-zero value in its records

Code Snippet `df.loc[:, ((df != 0).any(axis=0))]`



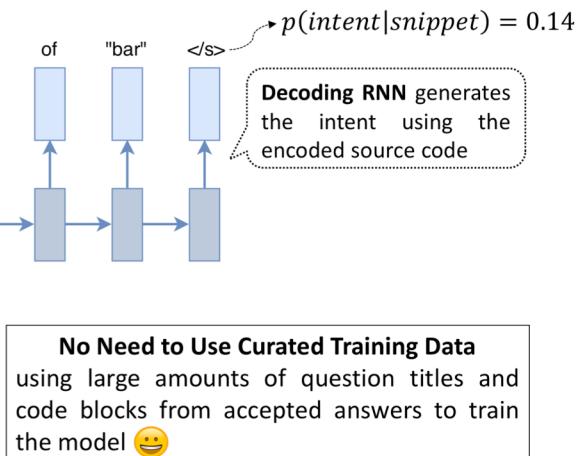
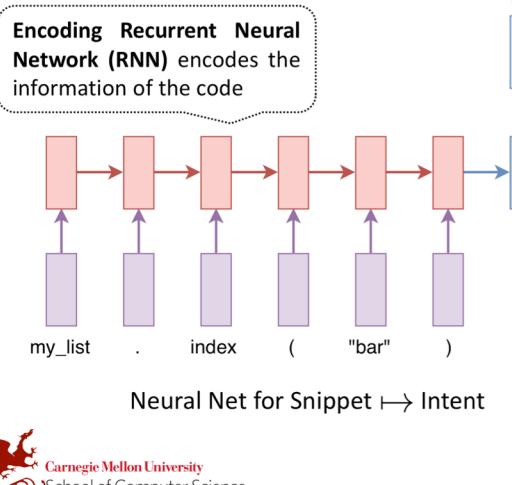
System Architecture

- Enumerate all possible candidate intent/snippet pairs in a Stack Overflow page
- Learn a classifier to rank each candidate intent/snippet pair



Neural Correspondence Model between Code and Intent

- **Neural sequence-to-sequence networks** for translation probability between intent and snippet



8

Examples

- Examples covers a wide variety of use cases
 - Built-in data type operation
 - I/O operation
 - Third-party library usage
- Examples are highly expressive and compositional!
 - Pose challenges to existing code/NL models

Intent dict how to create key or append an element to key
Rewritten Intent Create a key `key` if it does not exist in dict `dic` and append element `value` to value

Code Snippet `dic.setdefault(key, []).append(value)`

Intent How do I check if all elements in a list are the same

Rewritten Intent check if all elements in list `mylist` are the same

Code Snippet `len(set(mylist)) == 1`

Intent Iterate through words of a file in Python

Rewritten Intent check if all elements in list `mylist` are the same

Code Snippet `words = open('myfile').read().split()`

Intent Delete Column in Pandas based on Condition

Rewritten Intent delete all columns in DataFrame `df` that do not hold a non-zero value in its records

Code Snippet `df.loc[:, ((df != 0).any(axis=0))]`

Check our Dataset at
conala-corpus.github.io

Annotate Gold-standard Dataset

- Our system needs a *small* set of gold-standard intent/snippet data to learn the parameters of the classifier powered by high-level features
- Annotate top-ranked *how-to* questions on Stack Overflow for each language

	<h3>Gold Annotations</h3>
	527 examples
	142 questions
	330 examples
	100 questions

The majority of answers explain how to find a **single index**, but their method indexes if the item is in the list multiple times. Use `enumerate()`:

436

```
for i, j in enumerate(['foo', 'bar', 'baz']):
    if j == 'bar':
        print(i)
```

The `index()` function only returns the first occurrence, while `enumerate()` returns all occurrences.

As a list comprehension:

```
[i for i, j in enumerate(['foo', 'bar', 'baz']) if j == 'bar']
```

Here's also another small solution with `itertools.count()` (which is probably more efficient than the approach as `enumerate()`):

```
from itertools import izip as zip, count # izip for maximum
[i for i, j in zip(count(), ['foo', 'bar', 'baz']) if j == 'bar']
```

This is more efficient for larger lists than using `enumerate()`:

```
$ python -m timeit -s "from itertools import izip as zip, count # izip for maximum" "for i, j in zip(count(), ['foo', 'bar', 'baz']): if j == 'bar': print(i)" 10000 loops, best of 3: 174 usec per loop
$ python -m timeit "[i for i, j in enumerate(['foo', 'bar', 'baz']): if j == 'bar': print(i)]" 10000 loops, best of 3: 196 usec per loop
```

share improve this answer edited Nov 1 '17 at 23:55 answered Jun 19 '13 at 22:31 TerryA 40.5k ● 8 ● 75 ● 97

Finding the index of an item given its intent from itertools import izip as zip(C)ontext [i for i, j in zip(count(), ['foo', 'bar', 'baz']) if j == 'foo']

Rewritten Intent

Save Not Sure Reset All

System Deployment

- We deploy our system on the top-50K Python-tagged questions on Stack Overflow, and collected ~600K ranked intent/code-snippet pairs

