

# Detecting Interpersonal Conflict in Issues and Code Review: Cross Pollinating Open- and Closed- Source Approaches



Huilian Sophie Qiu, Bogdan Vasilescu, Christian Kaestner, Carolyn Egelman, Ciera Jaspan, and Emerson Murphy-Hill

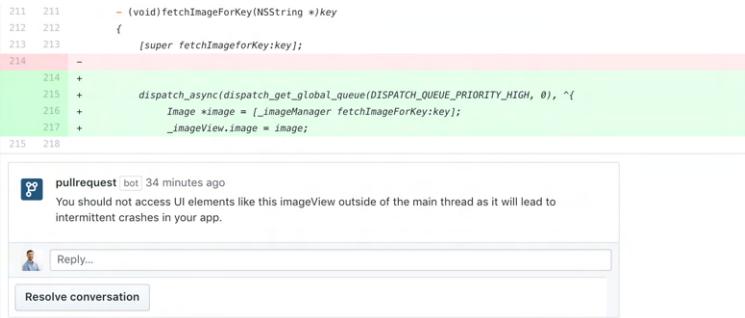
# Software development involves communication

## Issue discussions

-  Adding <a> Tag to redirect to Issue Link  
#9 opened 12 hours ago by alexmej1
-  updated searchbar tab  
#8 by Rnayak1 was merged 2 hours ago
-  Change color of active filter #3  
#7 by Yathousen was merged 12 hours ago
-  Github API Pagination good first issue hacktoberfest  
#6 opened 12 hours ago by billionverma
-  CSS Flexbox Fix good first issue hacktoberfest  
#5 by billionverma was closed 2 hours ago
-  Refactor css to scss good first issue hacktoberfest  
#4 by billionverma was closed 15 minutes ago

<https://github.com/dillionverma/github-issues-explorer/issues/17>

## Code reviews



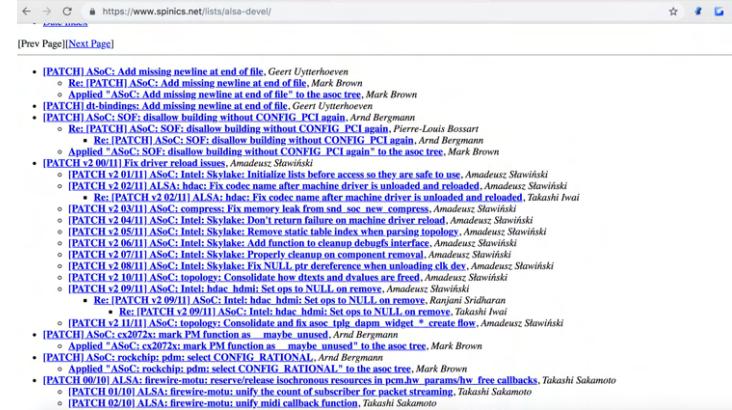
```
211 211 - (void)fetchImageForKey:(NSString *)key
212 212 {
213 213     [super fetchImageForKey:key];
214 214     -
215 215     dispatch_async(dispatch_get_global_queue(DISPATCH_QUEUE_PRIORITY_HIGH, 0), ^{
216 216         Image *image = [_imageManager fetchImageForKey:key];
217 217         _imageView.image = image;
218 218 }
```

 pullrequest bot 34 minutes ago  
You should not access UI elements like this imageView outside of the main thread as it will lead to intermittent crashes in your app.

 Reply...  
Resolve conversation

<https://www.pullrequest.com/blog/github-code-review-service/>

## Mailing lists



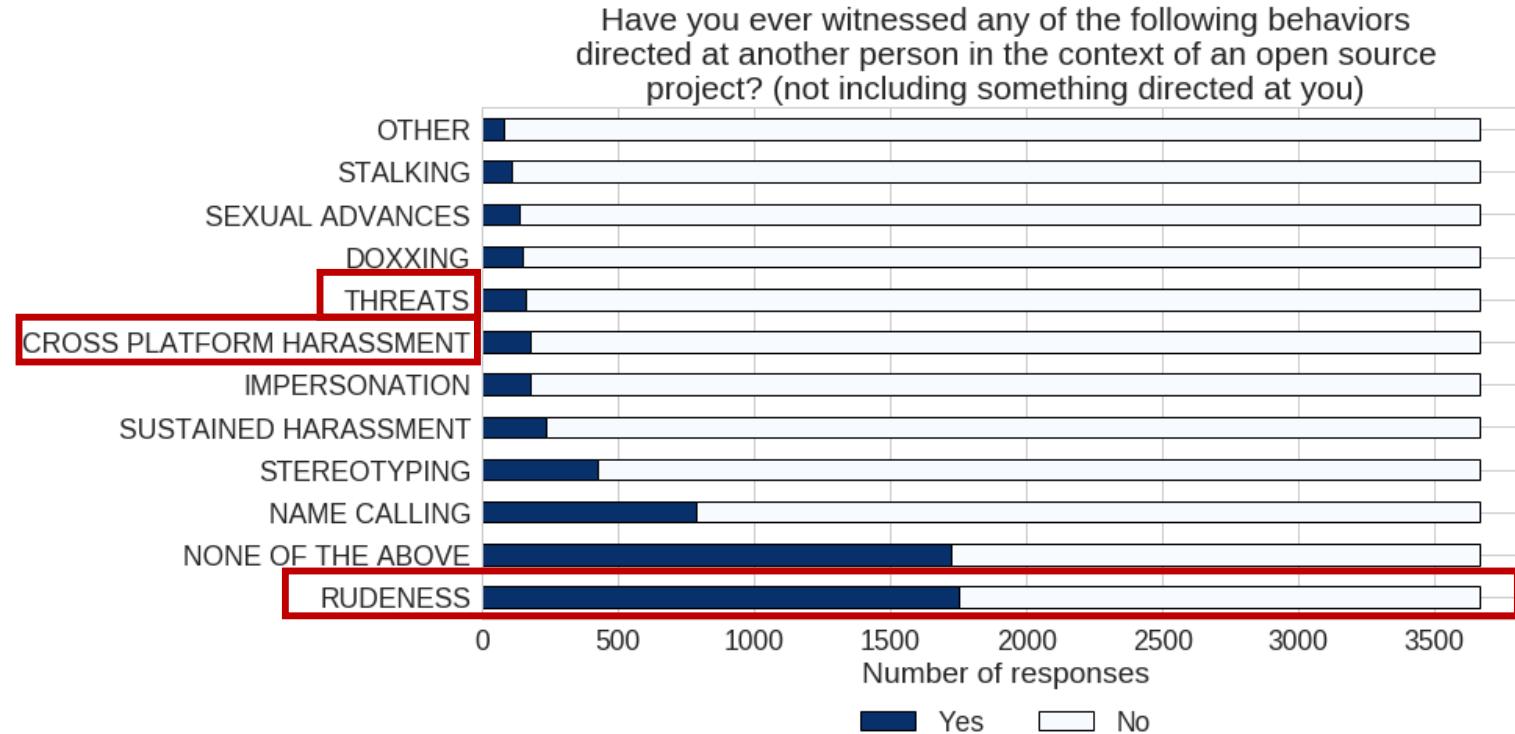
- [PATCH] ASoC: Add missing newline at end of file. Geert Uytterhoeven
- Applied "ASoC: Add missing newline at end of file." to the asoc tree. Mark Brown
- [PATCH] dt-bindings: Add missing newline at end of file. Geert Uytterhoeven
- [PATCH] ASoC: SOF: disable building without CONFIG\_PCI again. Arnd Bergmann
  - Re: [PATCH] ASoC: SOF: disable building without CONFIG\_PCI again. Pierre-Louis Bossart
  - [PATCH] ASoC: SOF: disable building without CONFIG\_PCI again. Arnd Bergmann
  - Applied "ASoC: SOF: disable building without CONFIG\_PCI again" to the asoc tree. Mark Brown
- [PATCH] v2.0(11) Fix driver reload issues. Andrasz Stawiński
  - [PATCH] v2.0(11) ASoC: Intel: Skylake: Initialize lists before access so they are safe to use. Andrasz Stawiński
  - [PATCH] v2.0(11) ASoC: Intel: Fix code flow after machine driver is unloaded and reloaded. Andrasz Stawiński
  - [PATCH] v2.0(11) ASoC: Intel: Fix code flow after machine driver is unloaded and reloaded. Takashi Iwai
  - [PATCH] v2.0(11) ASoC: Intel: Fix memory leak from and use after free. Andrasz Stawiński
  - [PATCH] v2.0(11) ASoC: Intel: Skylake: Don't return failure on machine driver reload. Andrasz Stawiński
  - [PATCH] v2.0(11) ASoC: Intel: Skylake: Remove static table index when parsing topology. Andrasz Stawiński
  - [PATCH] v2.0(11) ASoC: Intel: Skylake: Add function to cleanup debugfs interface. Andrasz Stawiński
  - [PATCH] v2.0(11) ASoC: Intel: Skylake: Fix memory leak from and use after free. Andrasz Stawiński
  - [PATCH] v2.0(11) ASoC: Intel: Skylake: Fix NULL pointer dereference when unbinding clk dev. Andrasz Stawiński
  - [PATCH] v2.0(11) ASoC: Intel: topology: Consolidate how dtsets and dtvalues are freed. Andrasz Stawiński
  - [PATCH] v2.0(11) ASoC: Intel: hdev: Set ops to NULL on remove. Andrasz Stawiński
    - Re: [PATCH] v2.0(11) ASoC: Intel: hdev: Set ops to NULL on remove. Ranjani Sriharan
    - Re: [PATCH] v2.0(11) ASoC: Intel: hdev: Set ops to NULL on remove. Takashi Iwai
- [PATCH] ASoC: ex2072x: mark PM function as \_\_maybe\_unused. Arnd Bergmann
- [PATCH] ASoC: rockchip: pmic: select CONFIG\_RATIONAL. Arnd Bergmann
  - Applied "ASoC: rockchip: pmic: select CONFIG\_RATIONAL" to the asoc tree. Mark Brown
- [PATCH] v3.0(10) ALSA: firewire-motu: resolve dependency between a pem-hv\_params/hv\_free callbacks. Takashi Sakamoto
  - [PATCH] v3.0(10) ALSA: firewire-motu: unify the count of subscriber for socket streaming. Takashi Sakamoto
  - [PATCH] v2.0(10) ALSA: firewire-motu: unify midi callback function. Takashi Sakamoto

<https://github.com/alice-sieve/Linux-Kernel/issues/72>

# However, these interactions are not always healthy

“FLOSS communities are such a **rich source of insulting or profane speech data.**”

(Squire et al., 2015)



(Geiger, 2017)

Some communities “neither reduce nor eliminate the **negative feedback**, as they **believe it is core** to the [code review] practice.”

(Alami et al., 2019)

# Negative communication experience can be discouraging



Negative reviews make one **begin to question  
one's confidence.**



**Done with contributing**



**... especially to marginalized groups**

(Nafus, 2012)

(Tourani et al., 2017)

# Automatic detection can help with prevention and mitigation

## Qualitative studies on the nature of negative interactions



Pushback (Egelman et al., 2020)



Incivility (Ferreira et al., 2021)



Interpersonal Conflicts (Gonçalves et al., 2022)



Toxicity (Miller et al., 2022)



Destructive criticism (Gunawardena et al., 2022)

Automatic tools can flag negative interactions and rephrase

Strategy plan	Input (upper) / Output (lower)
<u>Please</u> , <u>Subjunctive</u> , <u>Gratitude</u>	could you then please make some contributions in some of your many areas of expertise? thanks.
<u>Greeting</u> , <u>Subjunctive</u> , <u>Adverb.Just</u> , <u>For.Me</u>	hi , could you then just make some contributions for me in some of your many areas of expertise ?
<u>Please</u>	can someone please explain why there's a coi tag on this article? it's not evident from the talk page.
<u>For.Me</u> , <u>Hedges</u>	can someone explain why there ' s a coi tag on this article for me ? it ' s not apparent from the talk page .

(Fu et al., 2020)

# Unhealthy communication has been broadly studied in other online communities



New York Times comments

**rosegao/  
toxic\_tweets**

Toxic Tweet Detection and Trend Analysis

3 Contributors    0 Issues    0 Stars    1 Fork



Twitter

**Conv...Kit**

Cornell Conversational Analysis Toolkit

Wikipedia; Stack Exchange

# ...but off-the-shelf tools don't work well in SE

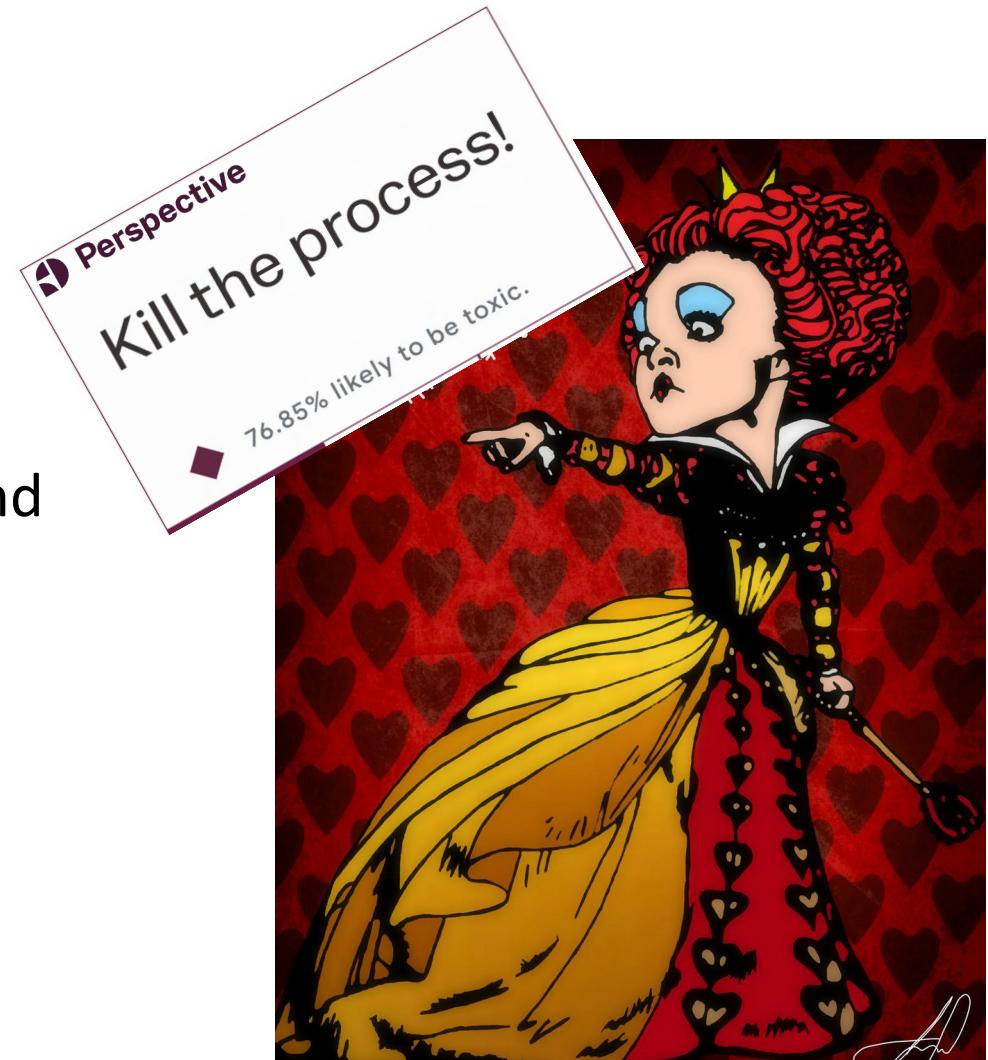
Off-the-shelf sentiment analysis tools have **low agreement** with software developers' labels.  
(Jongeling et al., 2015)

SE communications often include **technical jargon, word contractions, emoticons, URLs, and code snippets.** (Ahmed et al., 2017)

For example, domain words such as '**support**', '**error**', or '**default**' have neutral polarity (Islam and Zibran, 2017)



Kill ≠ Execute



# Several tools developed independently in SE



Sentiment analysis for SE (Islam and Zibran, 2017): JIRA issue comments



Sentiment analysis for SE (Ahmed et al., 2017): Gerrit code review conversations



Anger (Gachechiladze et al., 2017): Apache issue reports



Toxicity (Raman et al., 2020): GitHub issue conversations



Pushback (Egelman et al., 2020): Google code review conversations

# We cross-apply two recent works



(Raman et al., 2020)

OSS  
Issue



Corporate  
Code review



(Egelman et al., 2020)

# Research questions: how can we cross-pollinate the approaches



OSS  
Issue



Corporate  
Code review



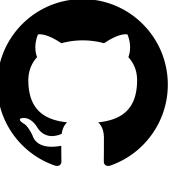
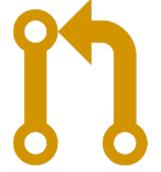
(Raman et al., 2020)

(Egelman et al., 2020)

# We cross-apply them along several axes



DETAILS IN  
THE PAPER

Types of communication	Concepts	Contexts	Method
 issues [Raman et al., 2020]	<b>Toxicity:</b> A rude, disrespectful, or unreasonable comment that is likely to make someone leave a discussion.		Text-based:  <b>Perspective</b>  <b>Politeness score</b>
 code reviews [Egelman et al., 2020]	<b>Pushback:</b> The perception of unnecessary interpersonal conflict in code review while a reviewer is blocking a change request.		Logs-based:  <b>Rounds of review</b>  <b>Active reviewing time:</b> time invested in providing feedback  <b>Active shepherding time:</b> time spent on actively viewing, responding to reviewer comments, or working on the selected code review

# Challenge 1: need to construct new datasets

## D1 Toxic OSS Issue Comments

[Raman et al., 2020]

Issues locked as **too-heated**

Contain the keyword “attitude”

173 toxic, 348 non-toxic



## D3 Pushback in Corporate Code Review Conversations [Egelman et al., 2020]

Survey programmers to label code reviews from a stratified random sample



Programmers report code reviews they perceive as pushback

493 pushback, 809 non-pushback



## D2 Toxic OSS Code Review Comments

122 toxic, 245 non-toxic



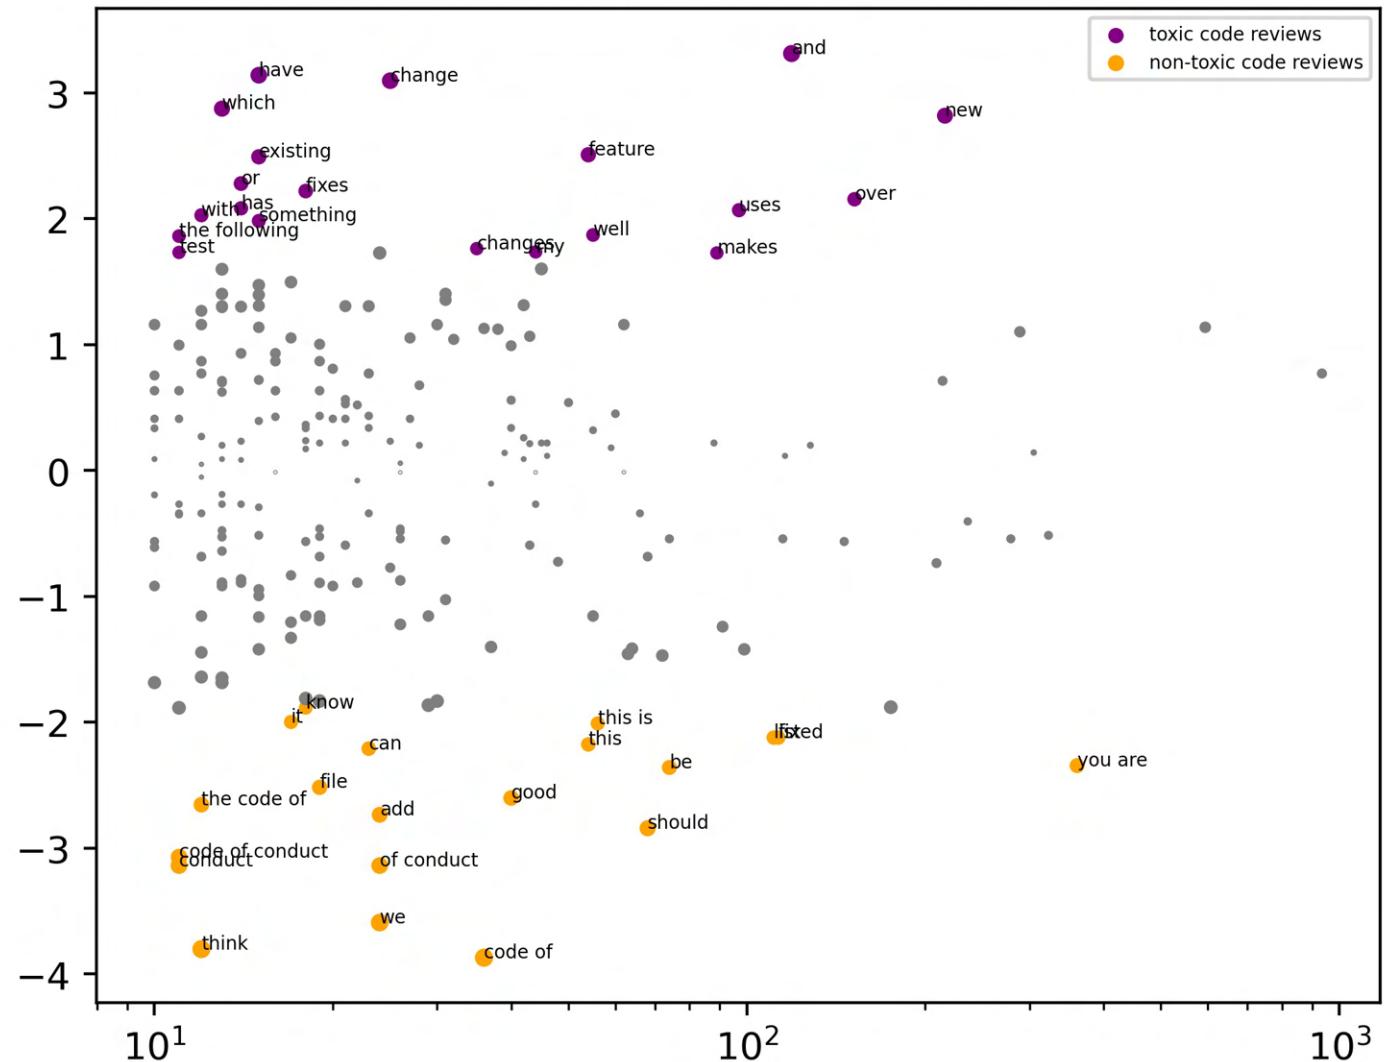
## D4 Pushback in OSS Code Review Conversations

228 pushback, 930 non-pushback

# Challenge 2: features for new classifiers

# Use Log Odds-Ratio to identify overrepresented n-grams in each label

$$\log\left(\frac{P(A)}{P(\sim A)}\right)$$



# We see more second-person pronouns among toxic comments

label	unigram	bigram	ngram
toxic OSS code review comments	you	it is	if you want
	people	you want	it is not
	even	that is	do you think
	what	you are	
	is	trying to	
	want	if you	
	your	to do	
	why	you think	
non-toxic OSS code review comments	tests	could you	
	unit	the pull	
	files	of files	
	for	we can	
	test	pull request	
	from	code to	
	line	instead of	the pull request

# We see more second-person pronouns among toxic comments

label	unigram	bigram	ngram
toxic OSS code review comments	<b>you</b>	it is	<b>if you want</b>
	people	<b>you want</b>	it is not
	even	that is	<b>do you think</b>
	what	<b>you are</b>	
	is	trying to	
	want	<b>if you</b>	
	<b>your</b>	to do	
	why	<b>you think</b>	
non-toxic OSS code review comments	tests	<b>could you</b>	
	unit	the pull	
	files	of files	
	for	we can	
	test	pull request	
	from	code to	
	line	instead of	<b>the pull request</b>

# Non-toxic code review comments contain more technical terms

label	unigram	bigram	ngram
toxic OSS code review comments	<b>you</b>	it is	<b>if you want</b>
	people	<b>you want</b>	it is not
	even	that is	<b>do you think</b>
	what	<b>you are</b>	
	is	trying to	
	want	<b>if you</b>	
	<b>your</b>	to do	
	why	<b>you think</b>	
non-toxic OSS code review comments	<u>tests</u>	<b>could you</b>	
	<u>unit</u>	<u>the pull</u>	
	<u>files</u>	<u>of files</u>	
	for	we can	
	<u>test</u>	<u>pull request</u>	
	from	<u>code to</u>	
	line	instead of	<u>the pull request</u>

# Both labels contain second-person pronouns

label	unigram	bigram	ngram
pushback corporate code review comments	<tech1>	<b>you want</b>	<b>you want to</b>
	tests	want to	on nov at pm
	<tech2>	of these	nov at pm
	our	of our	
	build	is to	
	libraries	if we	
	break	depend on	
	thing	we use	
	submit	to represent	make sure the
non-pushback corporate code review comments	groups	to me	to do the
	feature	how about	not sure if
	<tech3>	to submit	seems to be
	thanks	<b>could you</b>	<b>do you have</b>
	section	for the	how do we
	the	change the	to change the
	for	thanks for	thanks for the

# Non-pushback comments contain more gratitude

label	unigram	bigram	ngram
pushback corporate code review comments	<tech1>	<b>you want</b>	<b>you want to</b>
	tests	want to	on nov at pm
	<tech2>	of these	nov at pm
	our	of our	
	build	is to	
	libraries	if we	
	break	depend on	
	thing	we use	
	submit	to represent	make sure the
	groups	to me	to do the
non-pushback corporate code review comments	feature	how about	not sure if
	<tech3>	to submit	seems to be
	<b>thanks</b>	<b>could you</b>	<b>do you have</b>
	section	for the	how do we
	the	change the	to change the
	for	<b>thanks for</b>	<b>thanks for the</b>

# Challenge 2: features for new classifiers

Text-based classifier: [Raman et al., 2020]



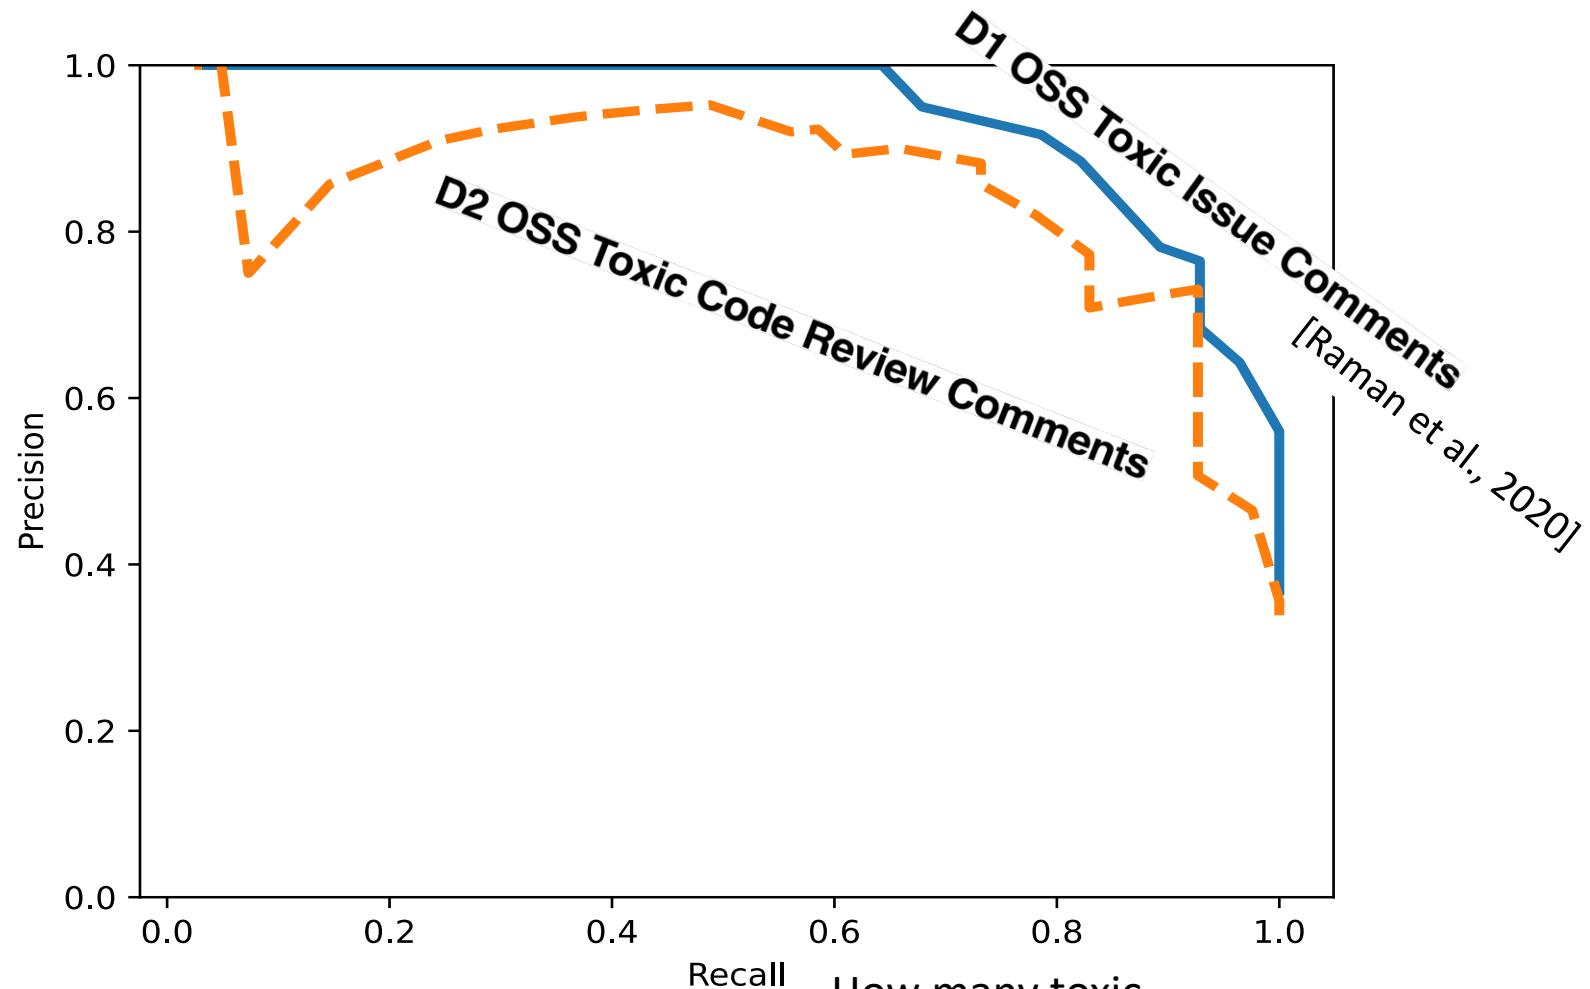
SentiCR (Ahmed et al., 2017)

Logs-based classifier: [Egelman et al., 2020]

- 
- The diagram illustrates the features used by a logs-based classifier. It shows three main components: 'Rounds of review' (represented by a speech bubble icon), 'Active reviewing time: time invested in providing feedback' (represented by a keyboard icon), and 'Active shepherding time: time spent on actively viewing, responding to reviewer comments, or working on the selected code review' (represented by a thumbs-up icon). A purple arrow points from the 'Rounds of review' icon to the 'Active reviewing time' icon, and another purple arrow points from the 'Active reviewing time' icon to the 'Active shepherding time' icon.
- Rounds of review
  - Active reviewing time: time invested in providing feedback
  - Active shepherding time: time spent on actively viewing, responding to reviewer comments, or working on the selected code review

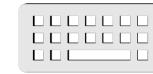
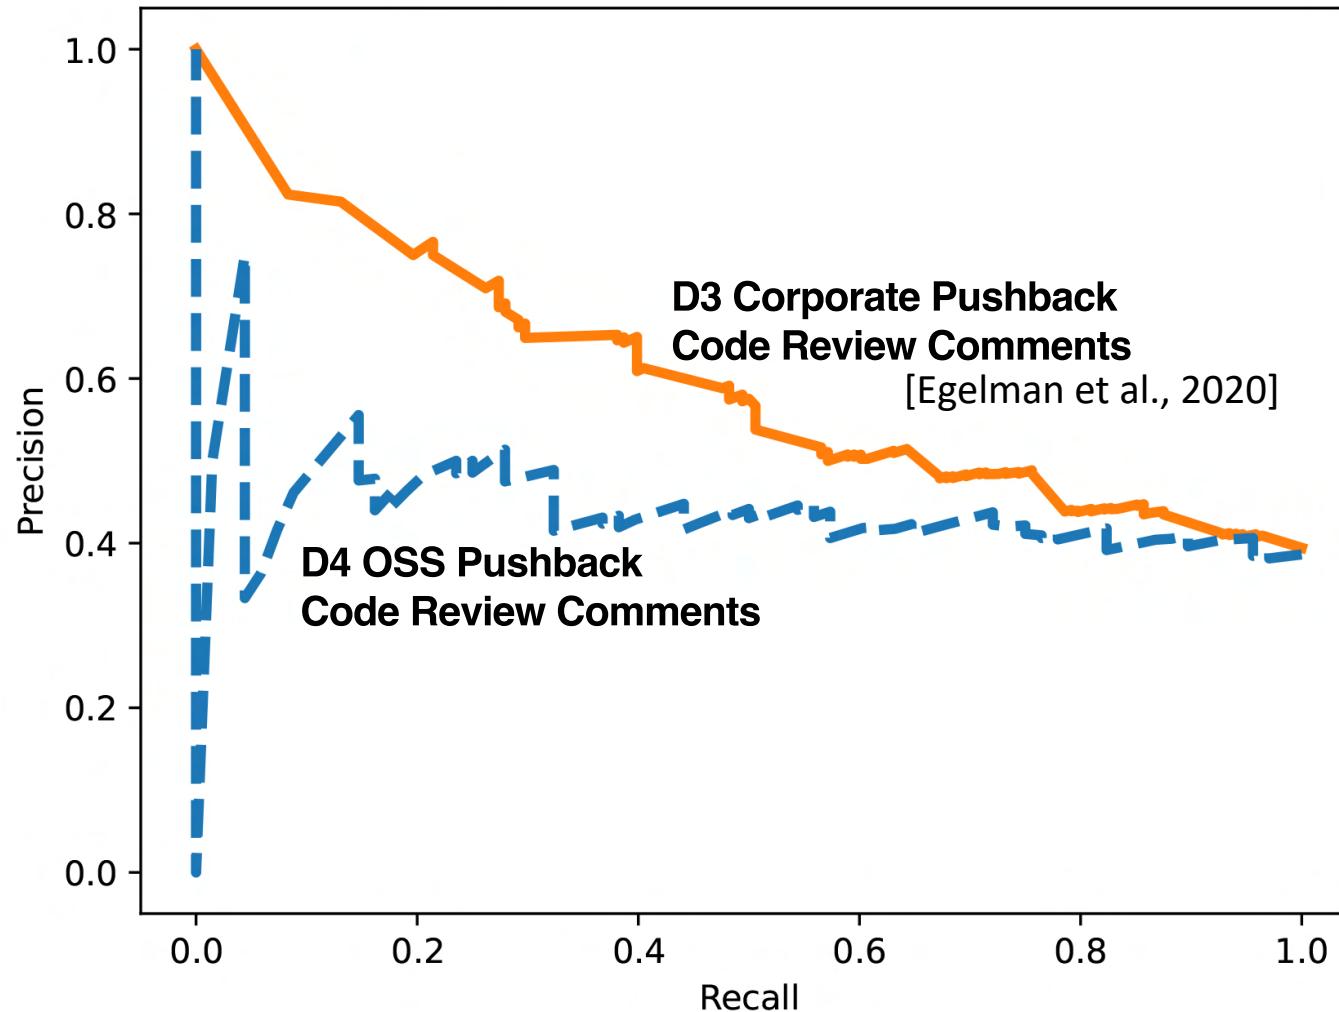
# Across types of communication: text-based classifier performs worse on code review comments

How many comments classify as toxic are really toxic



How many toxic comments have we classified as toxic

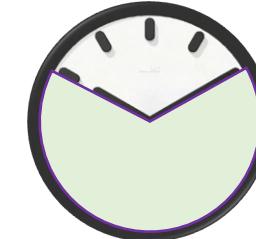
# Across contexts: logs-based classifier performs worse on open-source code review conversations



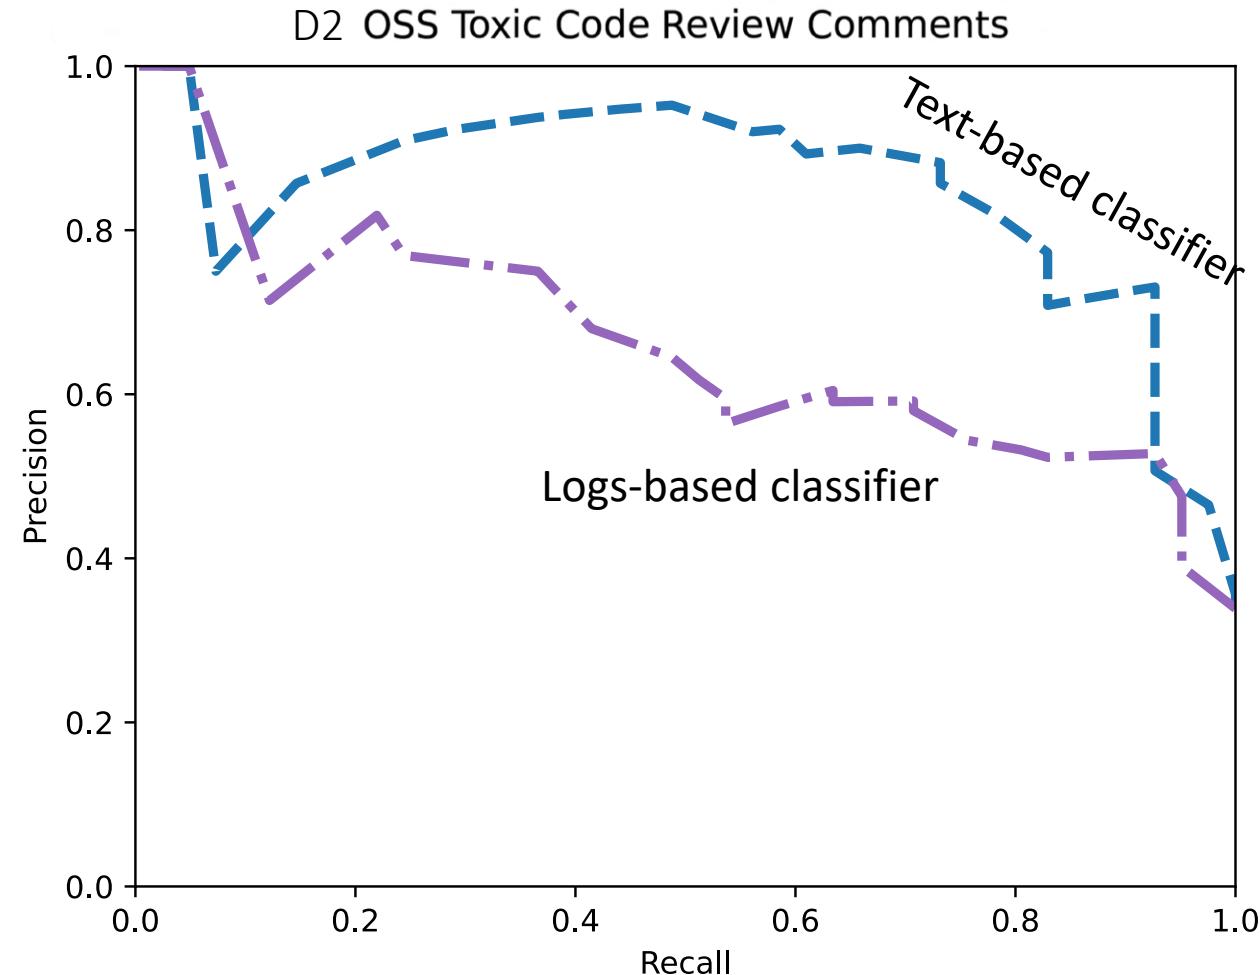
Active reviewing time: time invested in providing feedback



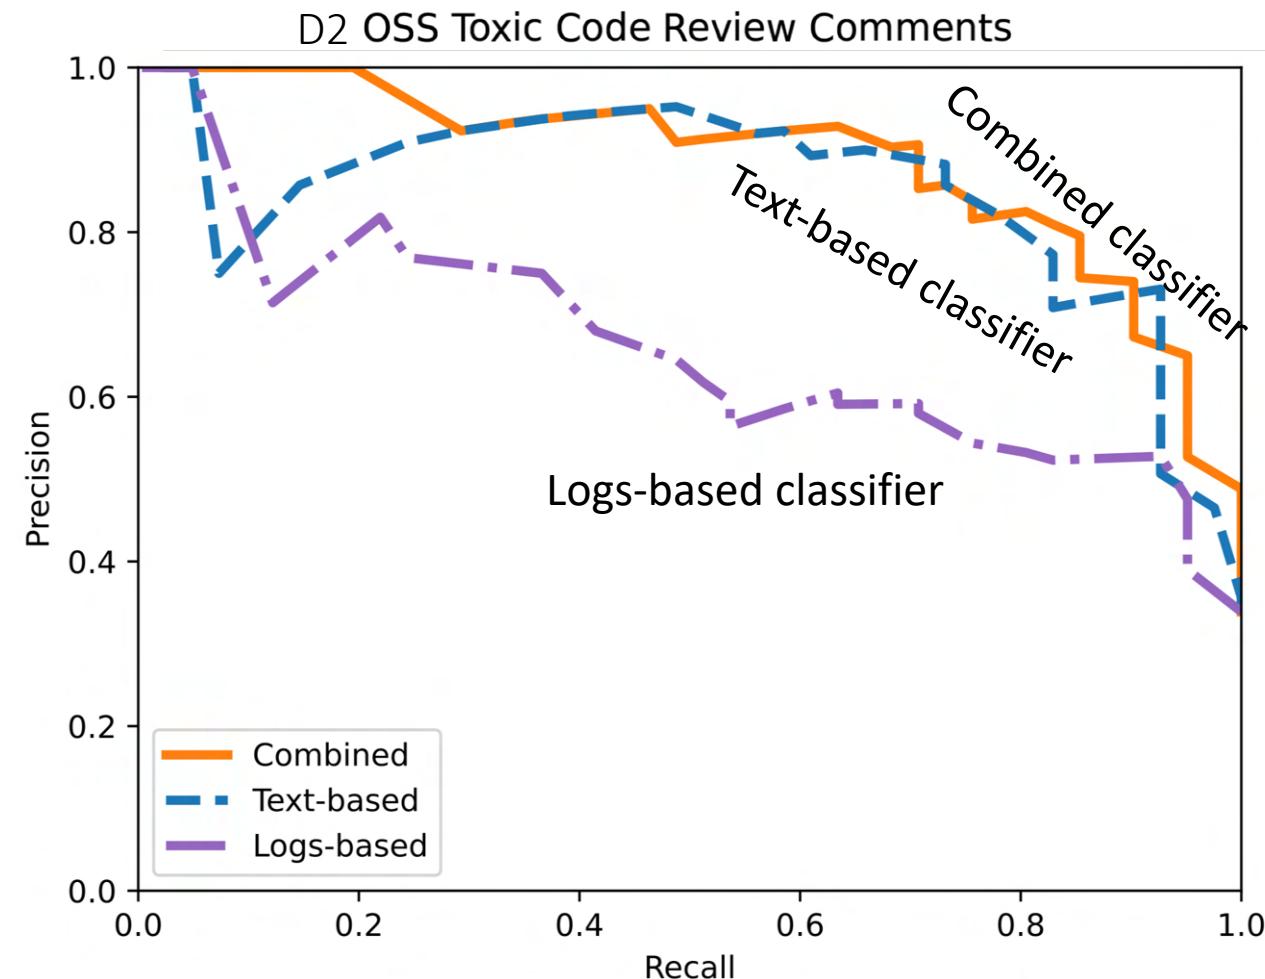
Active shepherding time: time spent on actively viewing, responding to reviewer comments, or working on the selected code review



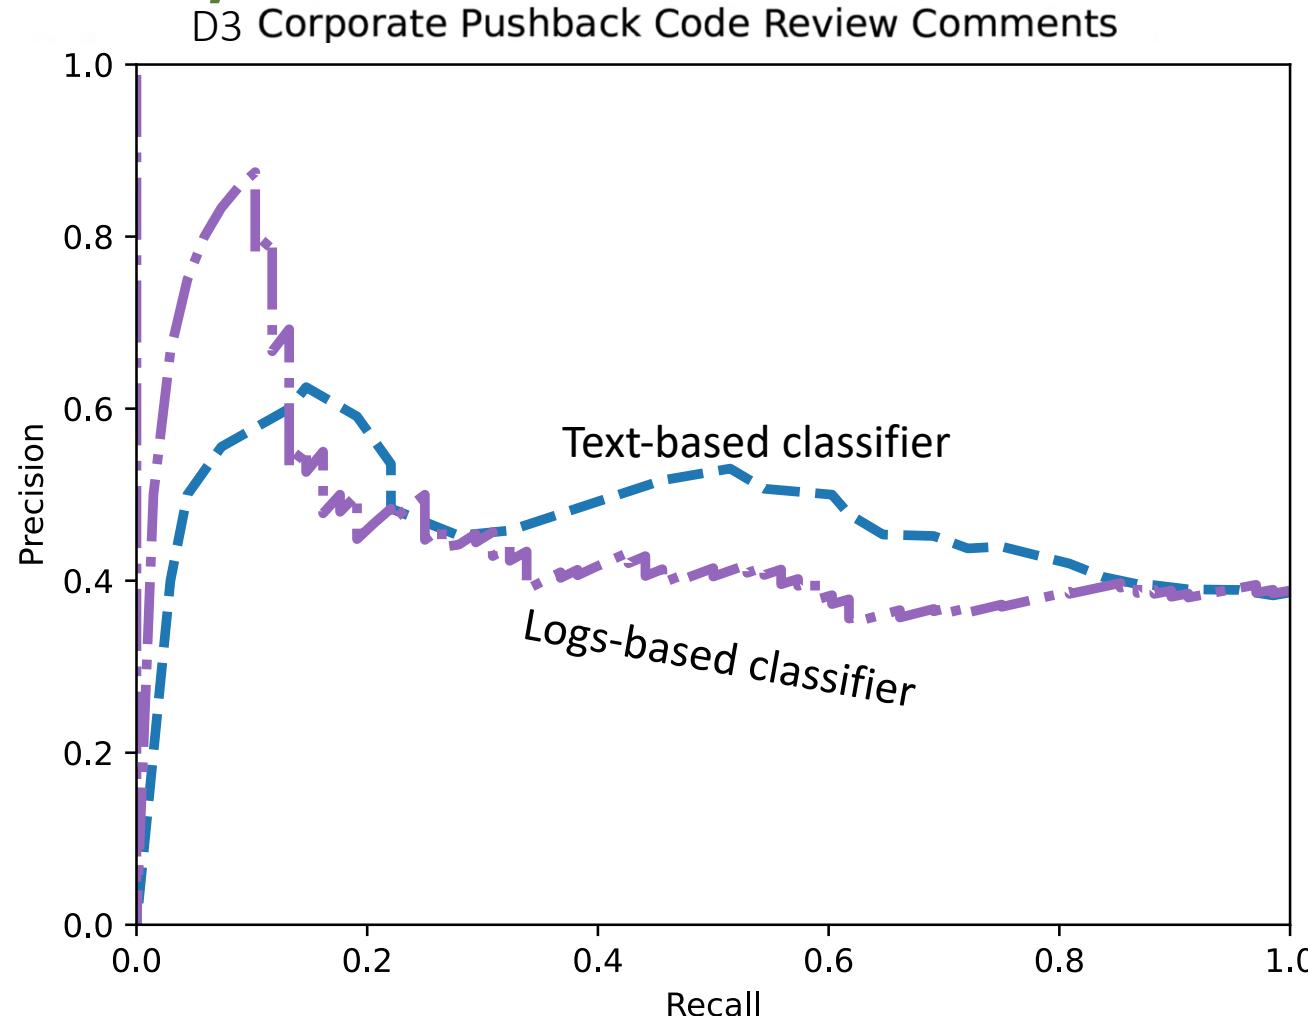
# Across methods: Toxicity is mainly about the use of inappropriate language



# Combined: Combined classifier performs better than the other two



# Across methods: Pushback in corporate has more nuance than delayed reviews



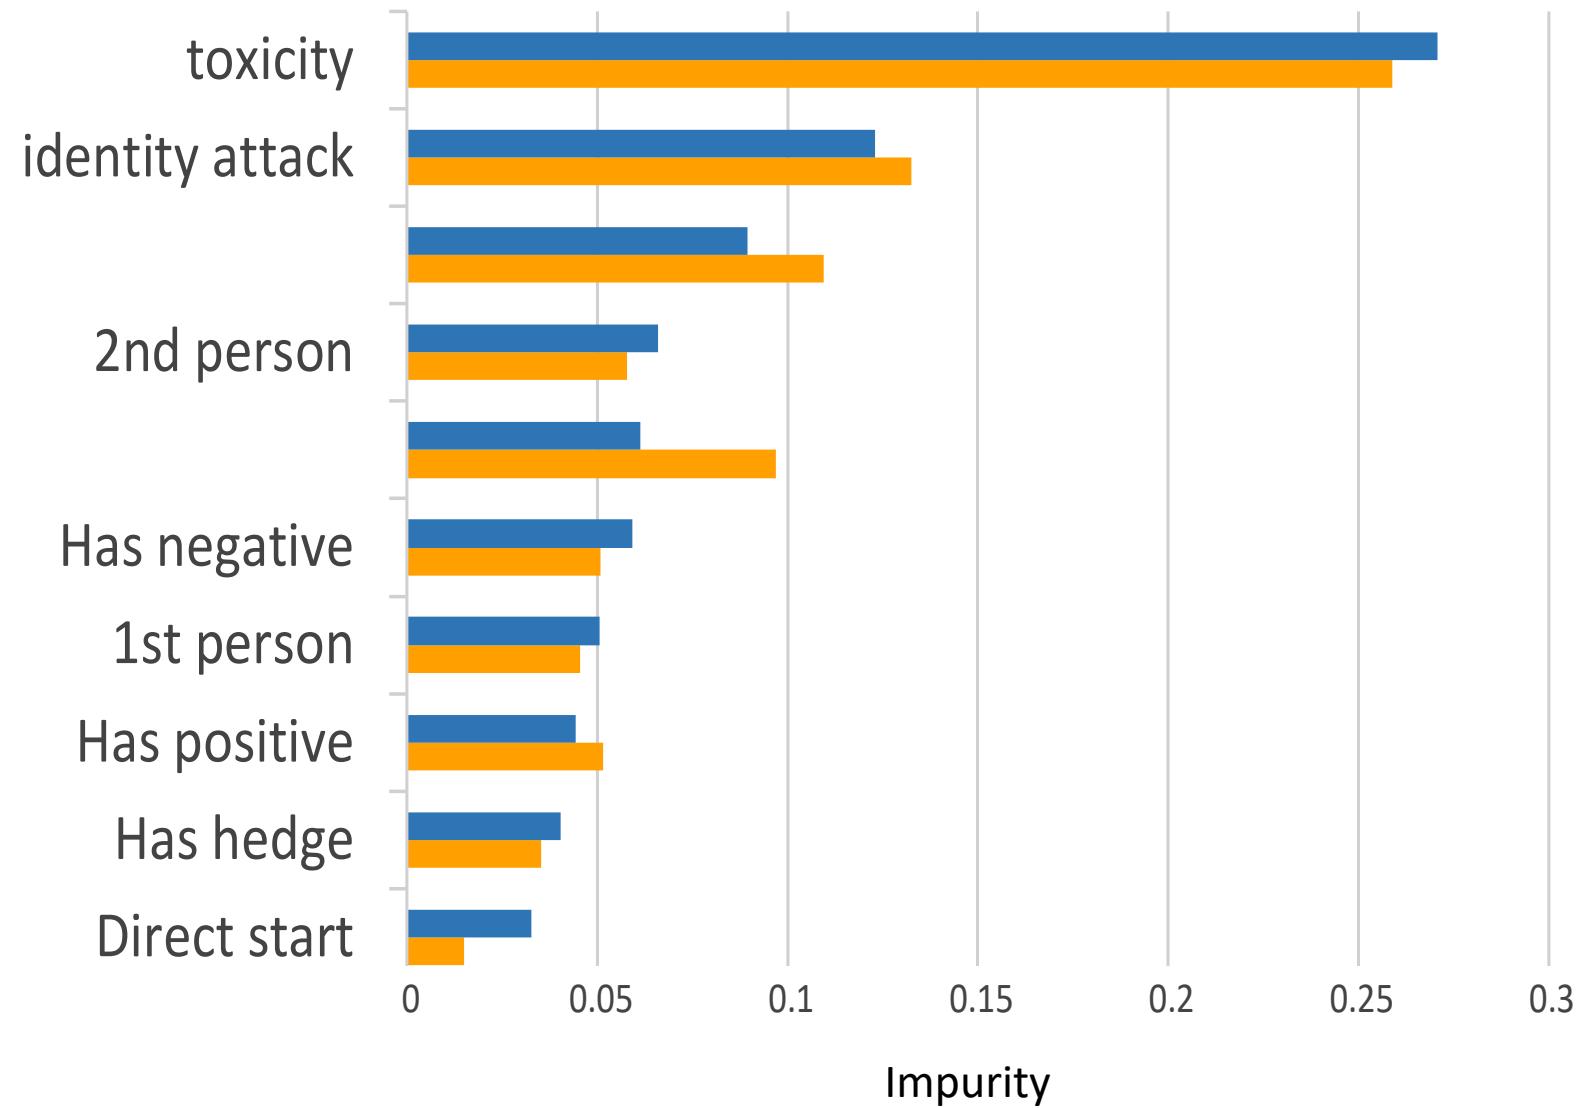
**Pushback:** The perception of unnecessary interpersonal conflict in code review while a reviewer is blocking a change request.

# Combined: Text-based features are important when detecting toxicity

D1 Toxic OSS Issue Comments

D2 Toxic OSS Code Review Comments

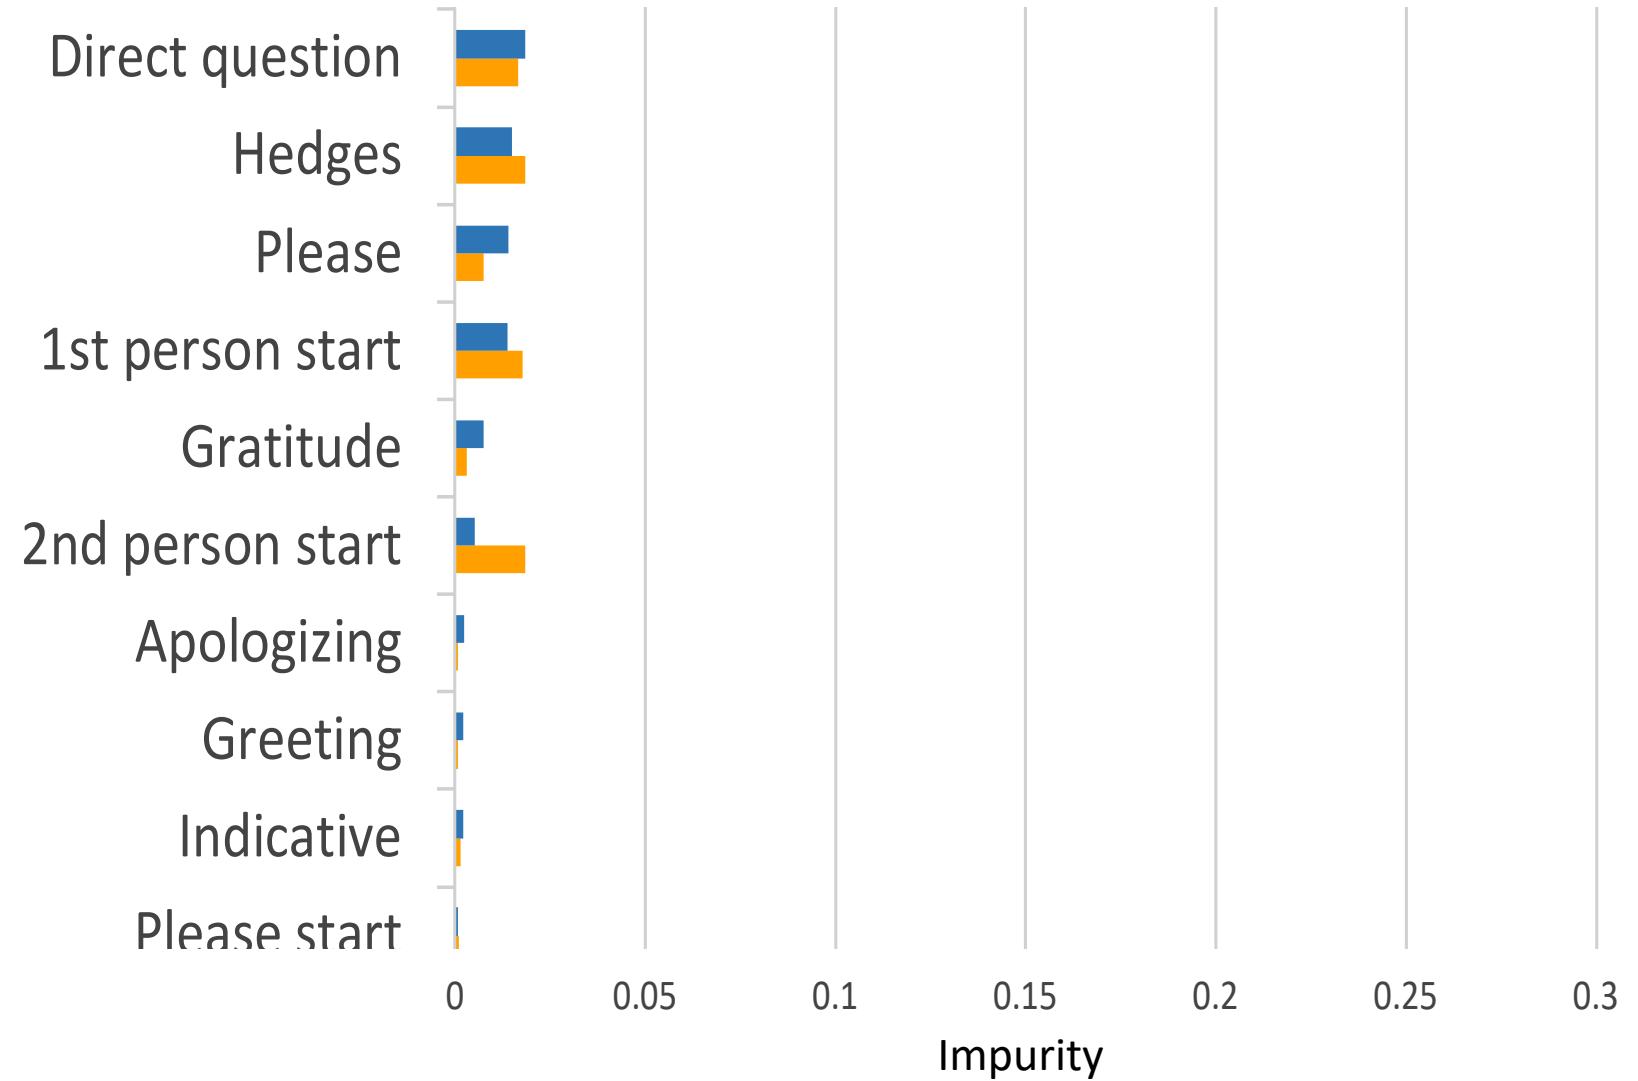
**Impurity: the probability of an item being incorrectly classified if it was randomly labeled according to the distribution of a specific feature.**



# Some politeness features are less important

D1 Toxic OSS Issue Comments

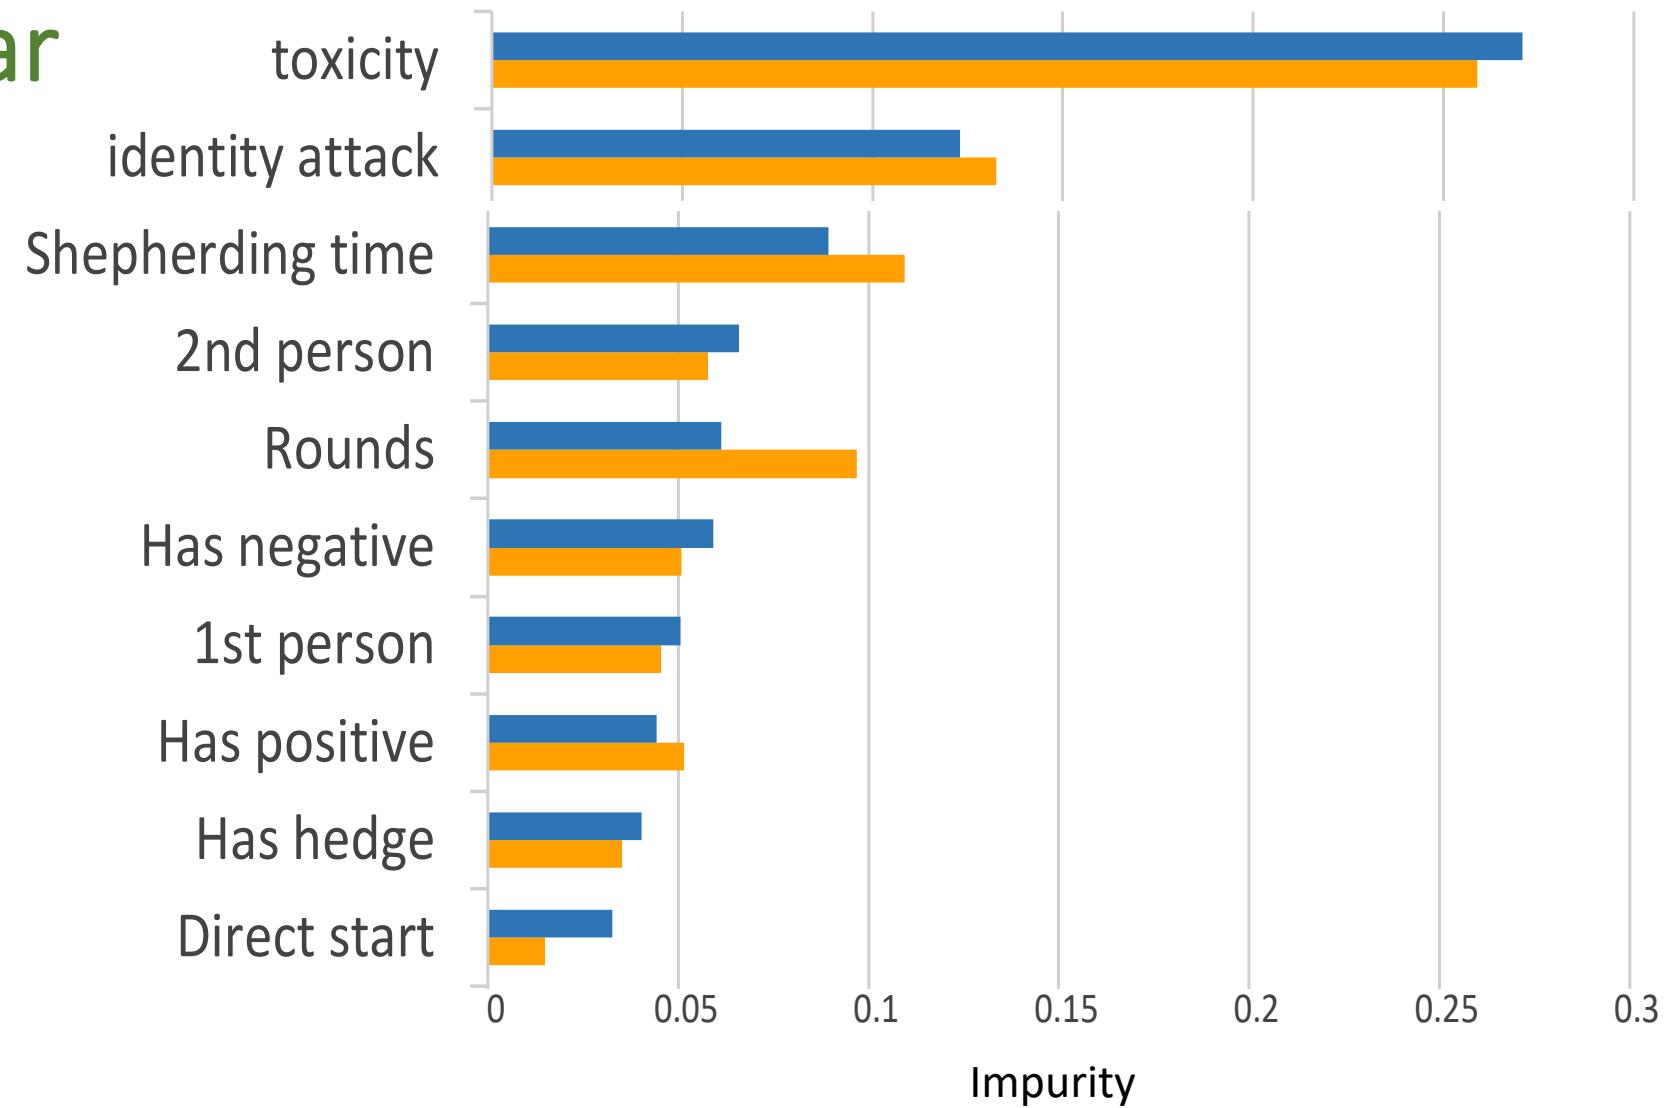
D2 Toxic OSS Code Review Comments



# Across types of communication: Feature importance rankings are similar

D1 Toxic OSS Issue Comments

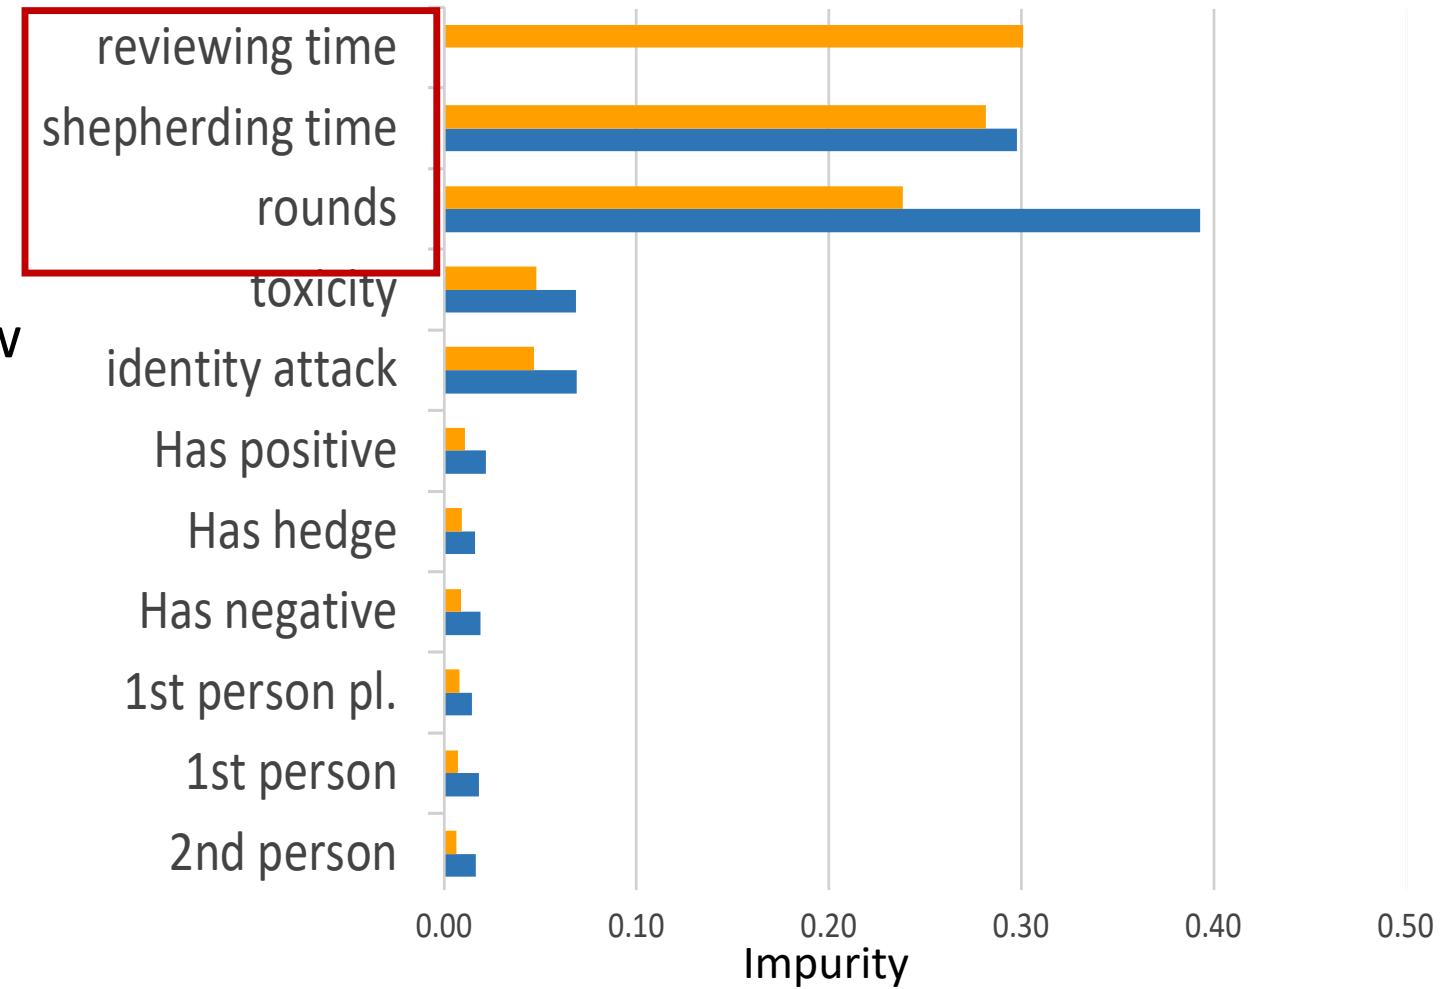
D2 Toxic OSS Code Review Comments



# Combined: Logs-based features are important for classifying pushback

D3 Corporate pushback  
code review

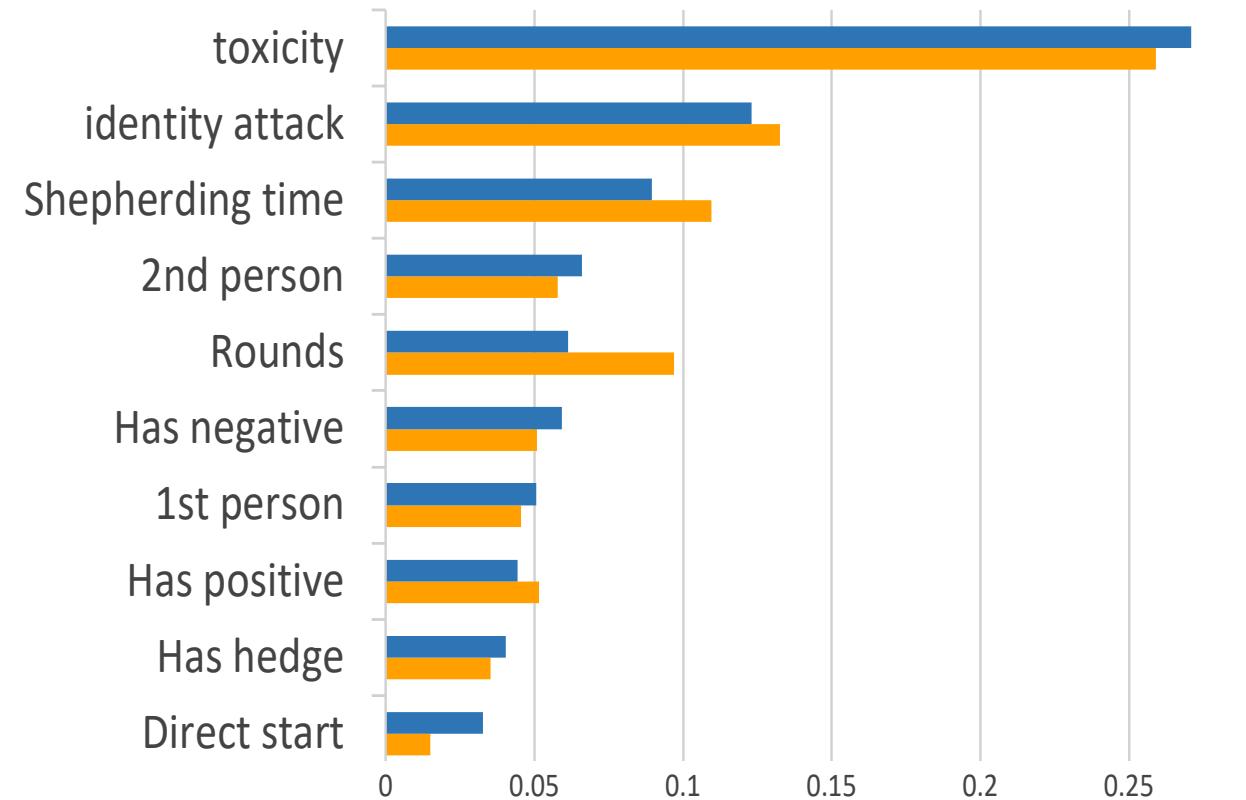
D4 OSS pushback code review



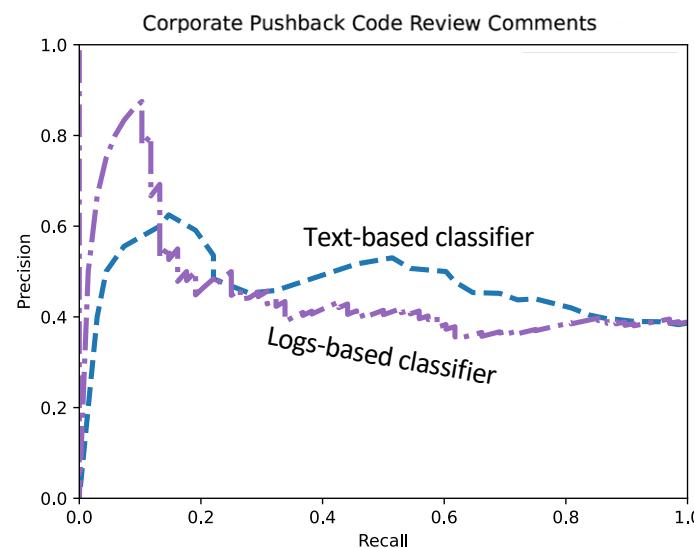
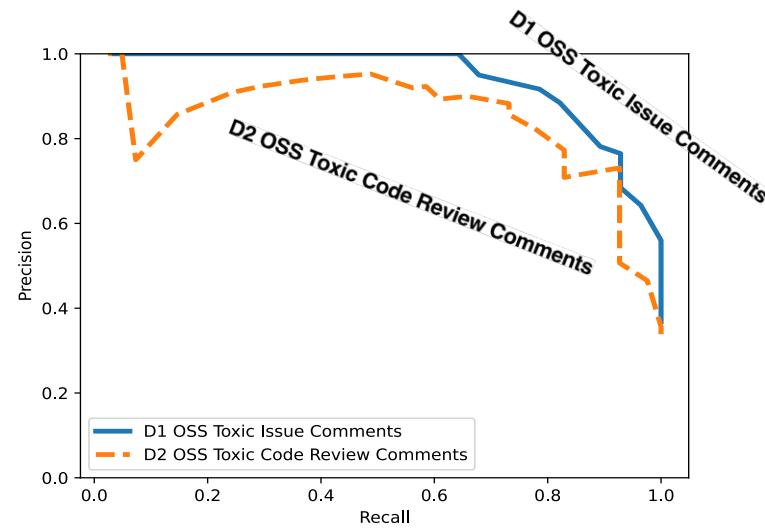
# Lots of potential to cross-pollinate

D1 Toxic OSS Issue  
Comments

D2 Toxic OSS Code Review  
Comments



# Open questions



## Improve the classifiers

More features

Deep Neural Networks



## Expand the datasets



## Prevention

# How can we improve our communication experience?



Comments in text can be interpreted very differently from how the intention was. **So always try to be extra polite when reviewing!**



label	unigram	bigram
toxic OSS code review comments	<b>you</b>	it is
	people	<b>you want</b>
	even	that is
	what	<b>you are</b>
	is	trying to
	want	<b>if you</b>
	<b>your</b>	to do
	why	<b>you think</b>

Use fewer 2<sup>nd</sup> person pronouns.  
Focus on the code.

label	unigram	bigram
non-pushback code review comments	<b>submit</b>	to represent
	groups	to me
	feature	how about
	<tech3>	to submit
	<b>thanks</b>	<b>could you</b>
	section	for the
	the	change the
	for	<b>thanks for</b>

Show more gratitude.

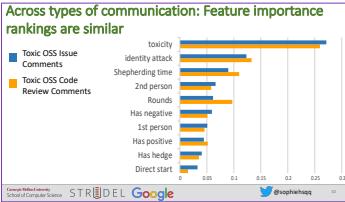
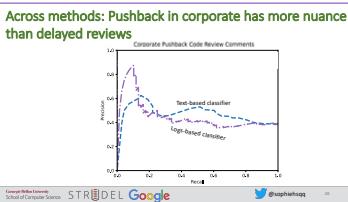
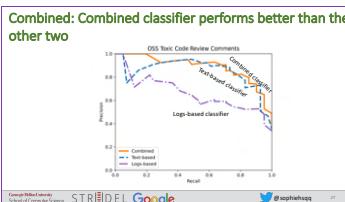
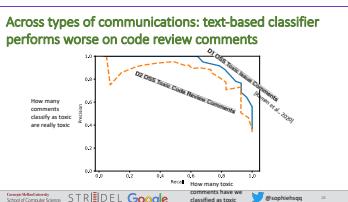
# Cross Pollinating Open- and Closed-Source Approaches

We cross-apply them along several axes 

Types of communication	Concepts	Contexts	Methods
 issues [Raman et al., 2020]	<b>Toxicity:</b> A rude, disrespectful, or unreasonable comment that is likely to make someone leave a discussion.		Text-based:  Perspective 
 code reviews [Egelman et al., 2020]	<b>Pushback:</b> The perception of unnecessary interpersonal conflict in code review while a reviewer is blocking a change request.		Logs-based:  Rounds of review  Active reviewing time: time invested in providing feedback  Active shepherding time: time spent on actively viewing, responding to reviewer comments, or working on the selected code review

DETAILS IN THE PAPER

## Result highlights



## Challenge 1: need to construct new datasets

### D1 Toxic OSS Issue Comments

[Raman et al., 2020]

Issues locked as too-heated  
Contain keywords "attitude"

173 toxic, 348 non-toxic



### D2 Toxic OSS Code Review Comments

122 toxic, 245 non-toxic



### D3 Pushback in Corporate Code Review Conversations

[Egelman et al., 2020]

Survey programmers to label code reviews from a stratified random sample



Programmers report code reviews they perceive as pushback

493 pushback, 809 non-pushback



### D4 Pushback in OSS Code Review Conversations

228 pushback, 930 non-pushback

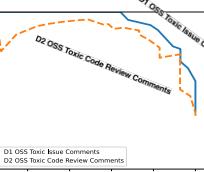


## Open questions

### Improve the classifiers

More features

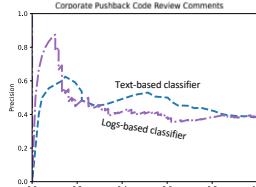
Deep Neural Networks



### Enlarge the datasets



amazon mechanical turk



### Prevention

